



Agent Architect Cohort 1

Agent Architecting

Day 2





Recap - Day 1 : Basic Concepts

- Understanding Large Language Model
- Different Types of LLMs
- Vector Databases
- Retrieval-Augmented Generation (RAG)
- Writing Effective Prompts
- How Agents actually evolved?
- LLM v/s Agents v/s Workflows
- Agent Orchestrations





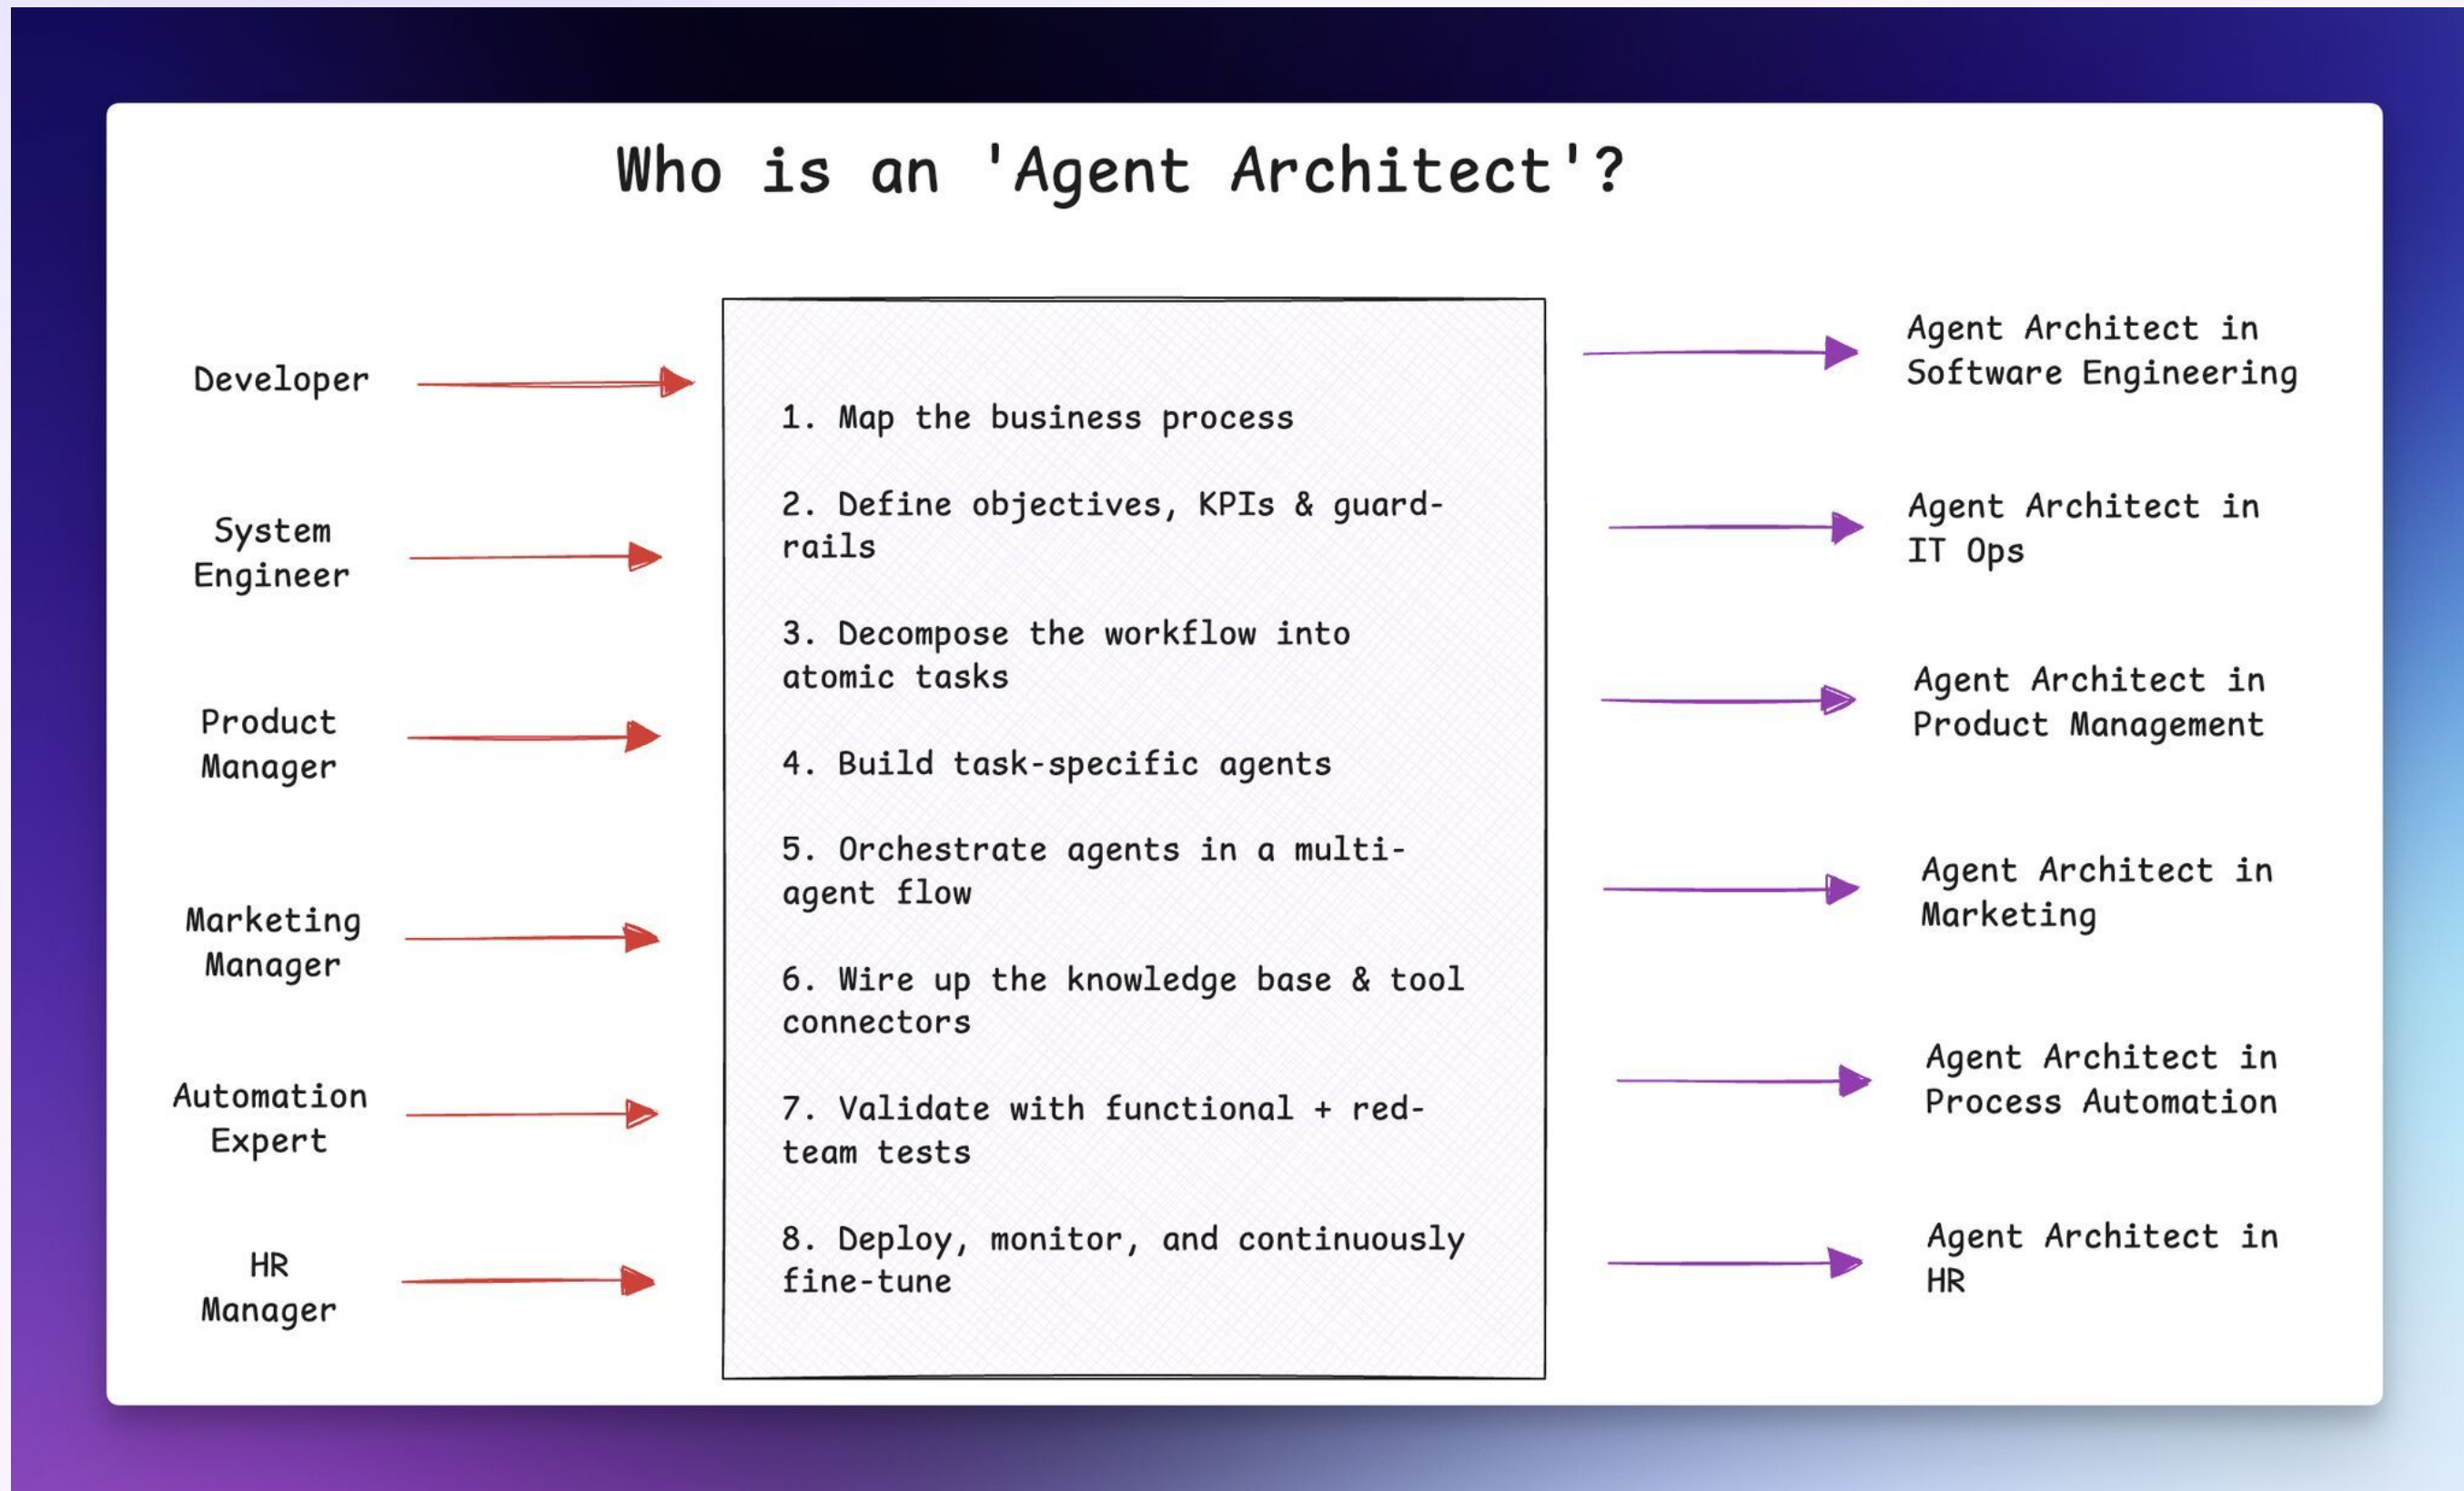
Agent Architect Cohort

Day 2 : Agent Architecting

- Business Requirements & AI Agents
- Core Components of an AI Agent (Tools, Functions, Extensions etc.)
- Agent Communication & MCP
- Safe & Responsible AI
- Secure Deployment of AI Agents
- Improving an AI agent
- Model Fine Tuning



Who is an 'Agent Architect'?



Business Requirements & AI Agents

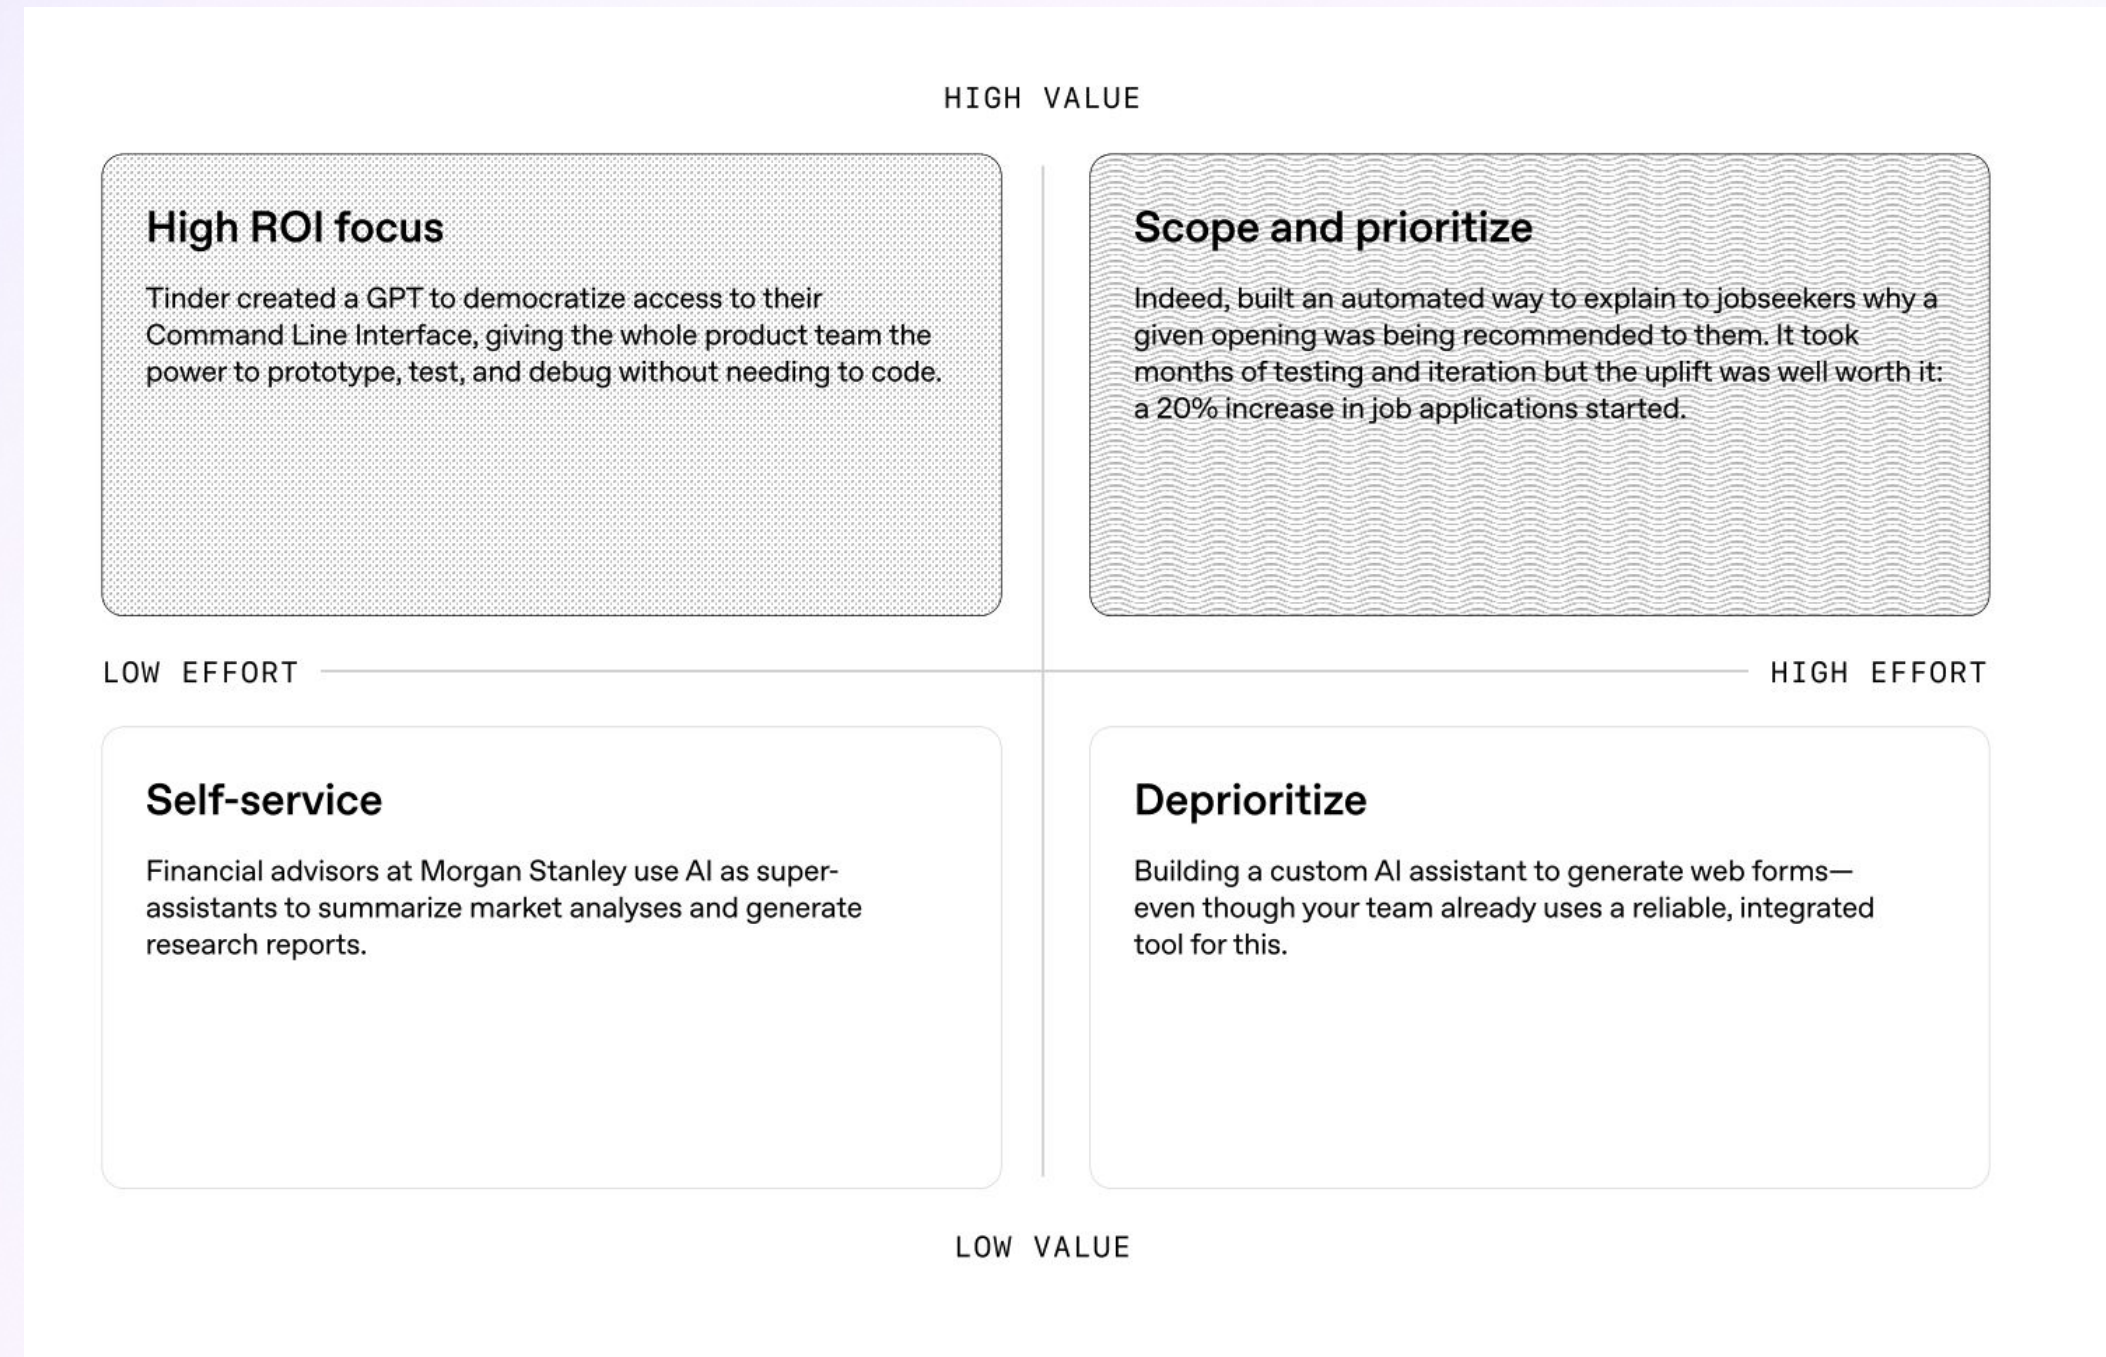
Before writing a single line of code or prompt, you must deeply understand the business context.

What problem are you trying to solve? Who are the users? What is the current process, and where are its pain points?

- **Identify Opportunities:** Look for areas within the business that can genuinely benefit from AI – tasks involving
 - complex decisions,
 - unstructured data, or
 - currently brittle rule-based systems.

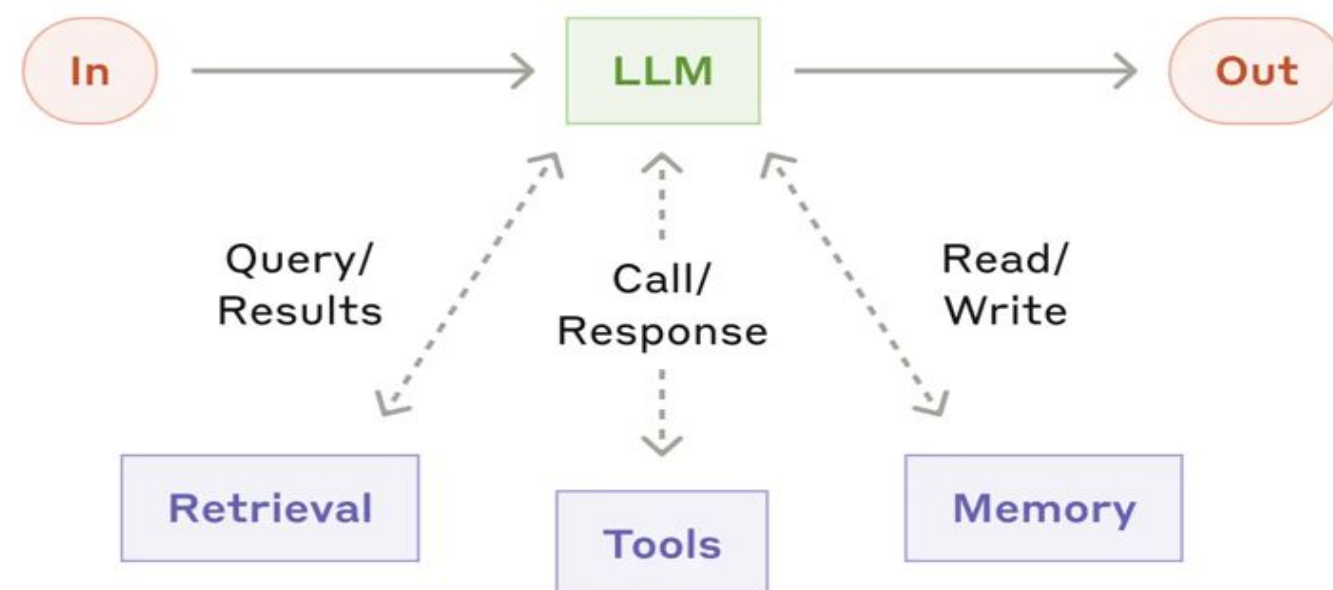
Focus on high-impact areas.

- **Stakeholder Collaboration:** Engage with business leaders, domain experts, and end-users to gather insights, define the scope, and understand the existing workflows.
- **Map the Current State:** Document the existing process. This helps identify inefficiencies and clarifies where an agent can provide the most value.



Core Components of an AI Agent : LLM

- The **Model** (the LLM itself) serves as the centralized decision-making unit for agent processes. It utilizes instruction-based reasoning and logic frameworks like **ReAct**, **Chain-of-Thought (CoT)**, or **Tree-of-Thoughts (ToT)** to plan and execute tasks.
- **The Augmented LLM** : The basic building block of agentic systems is an LLM enhanced with augmentations such as **retrieval, tools, and memory**. Current models can actively use these capabilities by generating their own search queries, selecting appropriate tools, and determining what information to retain.



Choosing the Right LLMs for Building Agents: Different models present varying strengths and trade-offs concerning task complexity, latency, and cost.

- **Capability vs. Task:** Match the model's capabilities (reasoning, tool use, multimodality) to the complexity of the tasks your agent needs to perform.
- **Accuracy First:** Start by prototyping with the most capable model available to establish a performance baseline and ensure your agent *can* achieve the desired accuracy.
- **Optimize for Cost & Latency:** Once accuracy targets are met, experiment with swapping in smaller, faster, or less expensive models for specific tasks or sub-agents, provided they maintain acceptable performance.
- **Context Window:** Consider the size of the context window required for your agent's tasks, especially if it needs to process large documents or maintain long conversations.

Core Components of an AI Agent : Tools

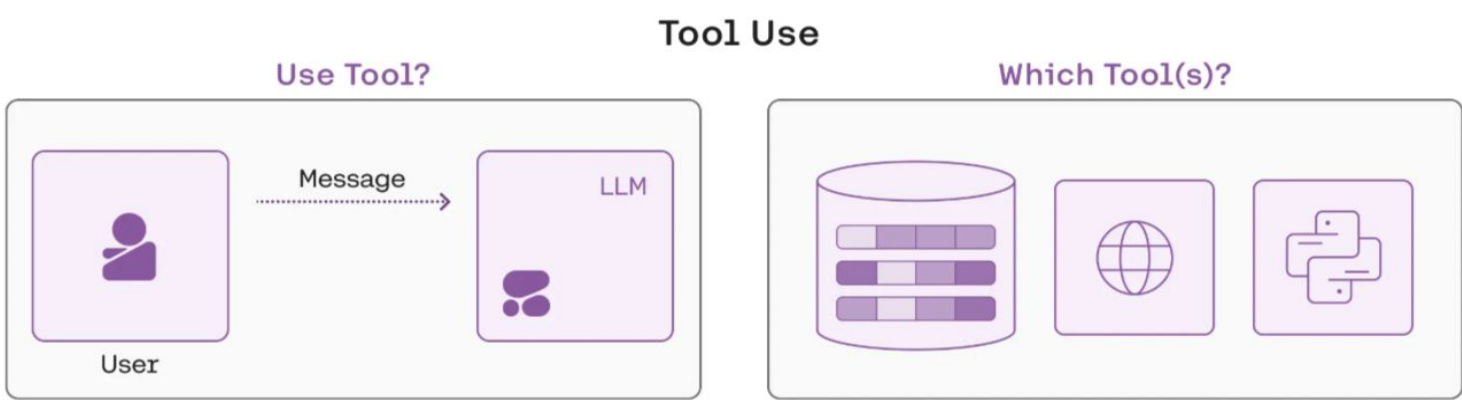
Tools extend agent capabilities beyond the LLM, enabling real-time data access, external actions, and system interactions.

Agents use tools to call APIs or interact with UIs, helping them gather context and perform meaningful tasks.

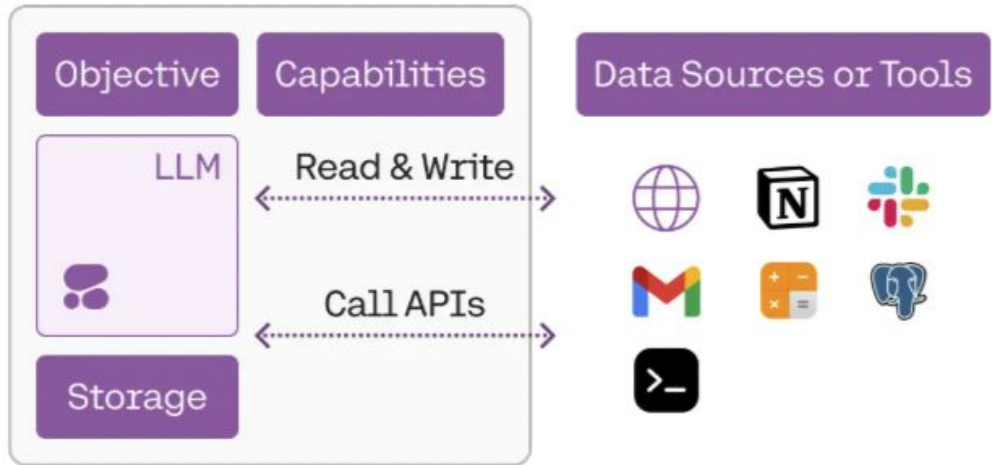
Standardized tools (with clear name, description, parameters, and outputs) ensure reusability, easier updates, and governance.

Tool types include:

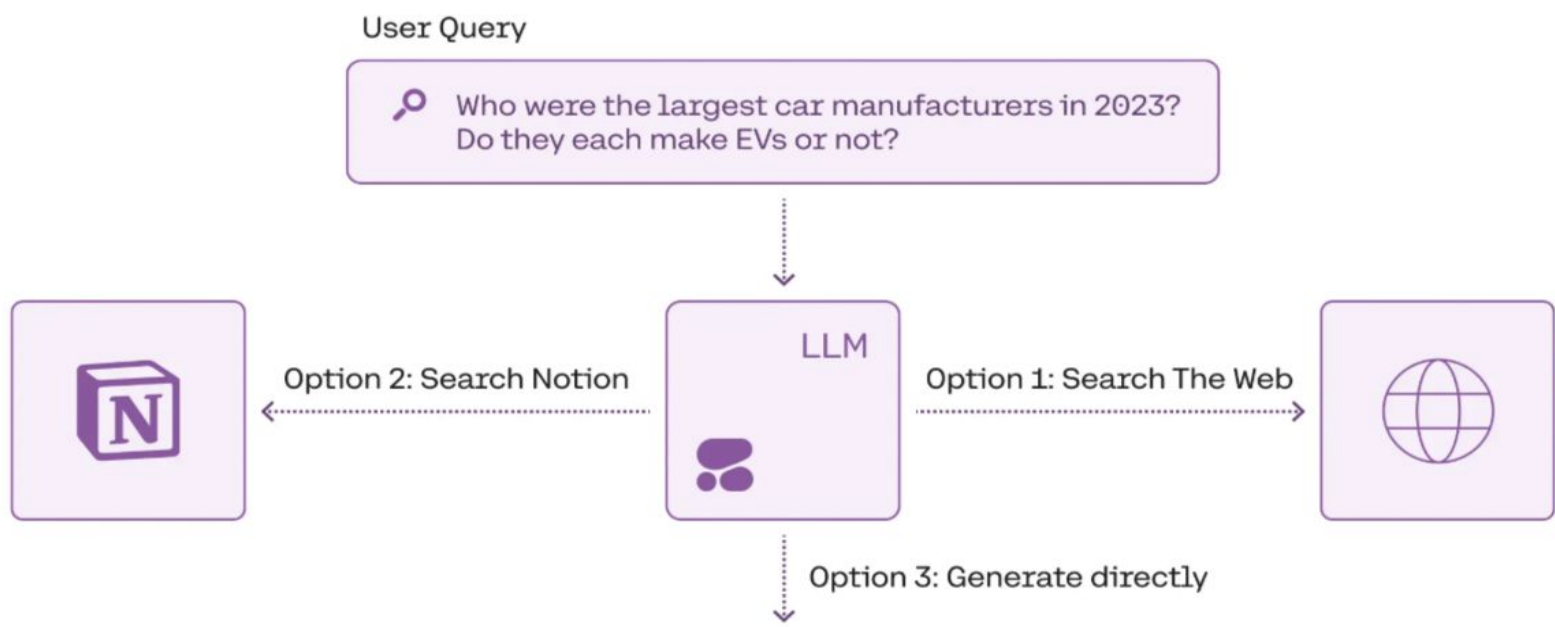
- *Data tools* (e.g., DB queries, web search)
- *Action tools* (e.g., send email, update CRM)
- *Agent tools* (agents acting as tools for orchestration)



Decision-making in tool use



Tool use enables LLMs to be more connected to the external world



Tool use enables routing a request to only the relevant tools

Deep Dive on Tools : Extensions

- Extensions **bridge the gap between an API and an agent** in a standardized way, allowing agents to seamlessly execute APIs. They teach the agent how to use an API endpoint (using examples) and what parameters are needed.
- *Sample Extension:* The **Code Interpreter** extension allows generating and running Python code from natural language descriptions.

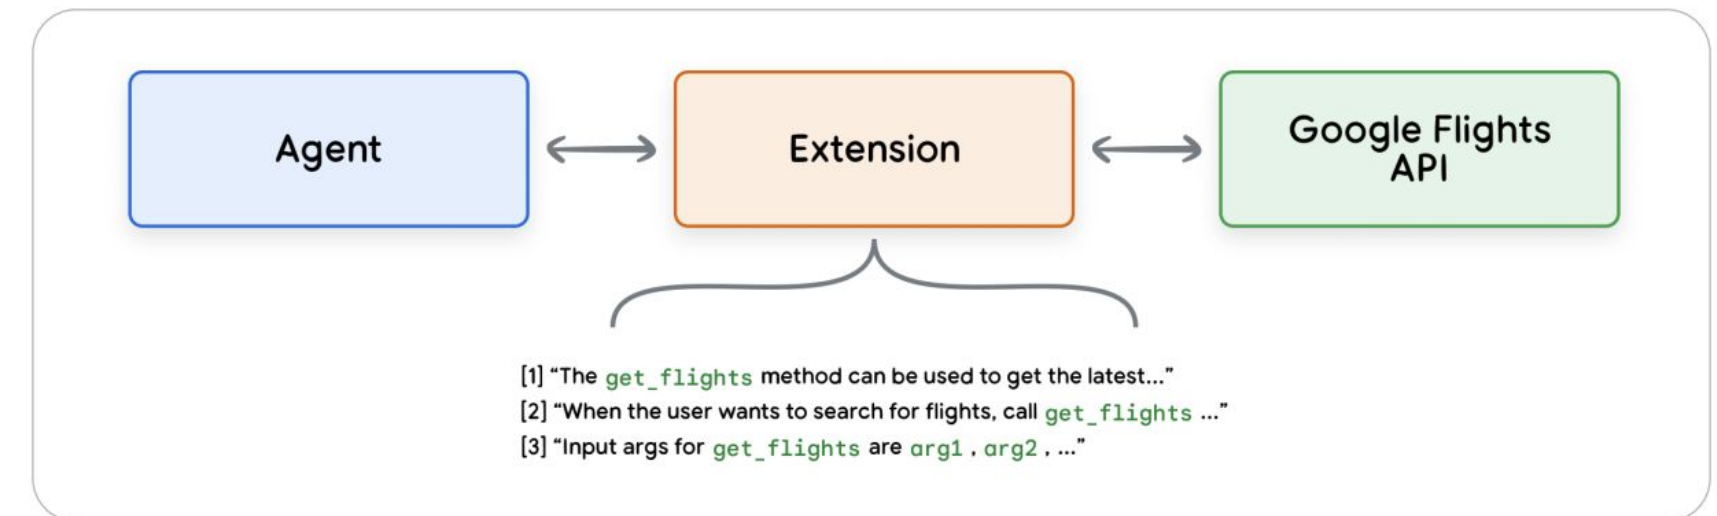


Figure 4. Extensions connect Agents to External APIs



Figure 5. 1-to-many relationship between Agents, Extensions and APIs

Deep Dive on Tools : Functions

- Functions are **self-contained modules of code** that accomplish specific tasks. In agent systems, the model decides which function to use and what arguments it needs based on its specification.
- Unlike extensions (which are typically executed agent-side and can make live API calls), functions are often structured so that the model outputs the function name and arguments, and the actual execution (e.g., the API call) happens on the client-side. **This gives developers more granular control.**
- *Use Cases:* Making API calls at different layers of the application stack, handling security/authentication restrictions, managing timing constraints (batch operations, human-in-the-loop), applying additional data transformations, or "stubbing" APIs during development.

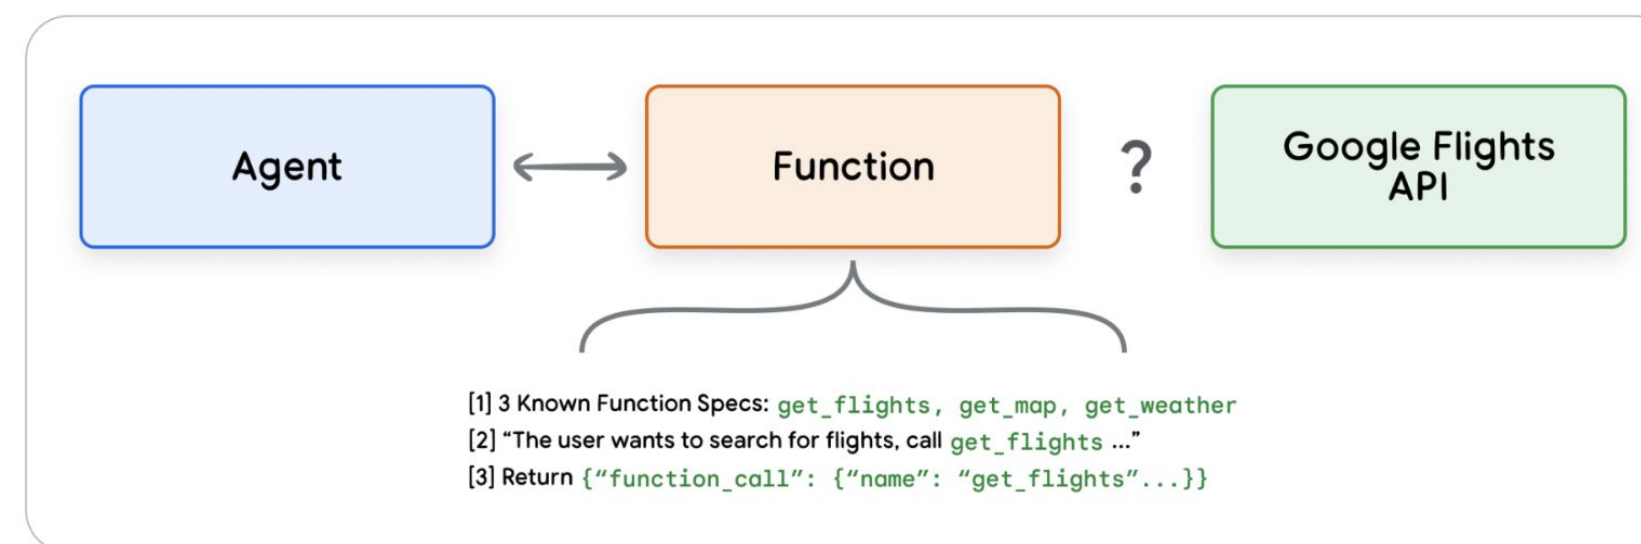


Figure 7. How do functions interact with external APIs?

Deep Dive on Tools : Data Stores

- Data Stores provide agents with access to dynamic, up-to-date, and often proprietary information beyond the model's training data, ensuring responses are grounded in factuality and relevance. This is commonly used in **Retrieval Augmented Generation (RAG)**.
- They allow developers to provide data in its original format (e.g., PDFs, spreadsheets, website content), which the Data Store converts into vector embeddings for the agent to use for information retrieval.

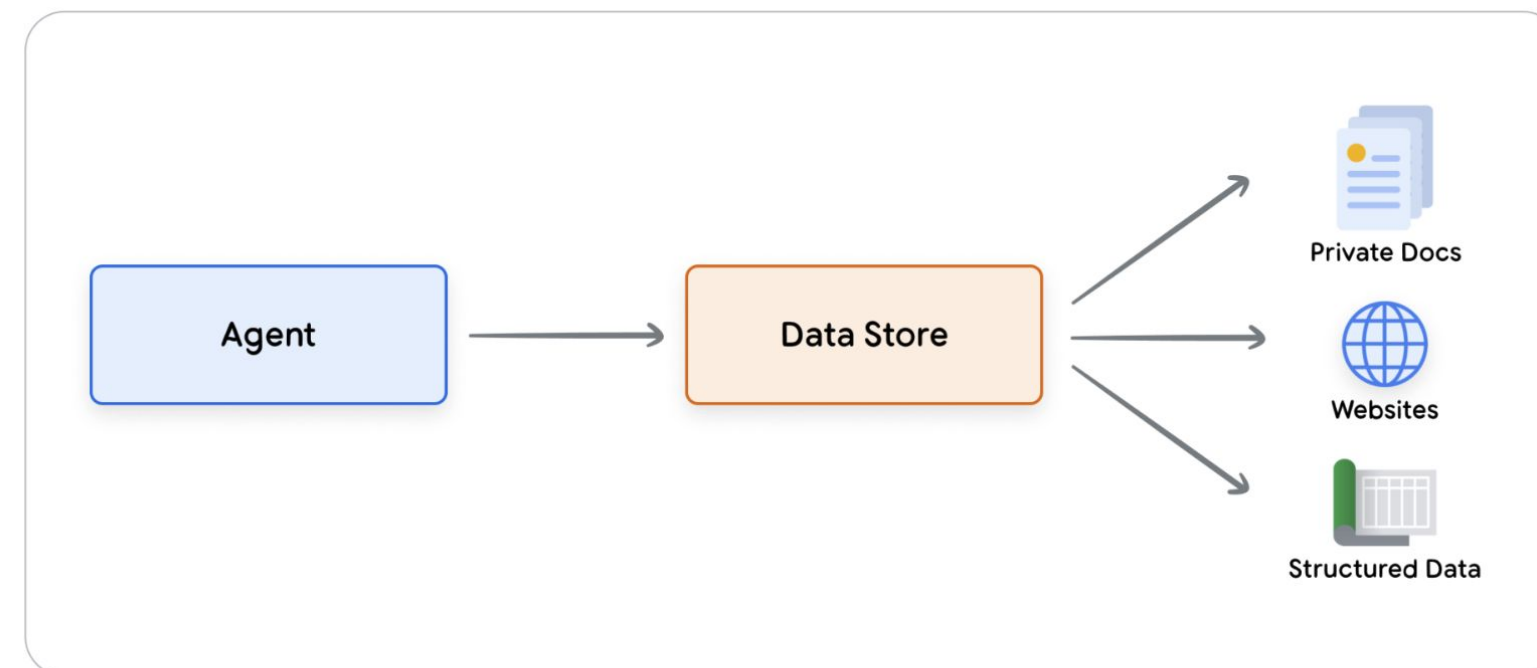


Figure 11. Data Stores connect Agents to new real-time data sources of various types.

Summarising Tools

| | Extensions | Function Calling | Data Stores |
|-----------|---|--|--|
| Execution | Agent-Side Execution | Client-Side Execution | Agent-Side Execution |
| Use Case | <ul style="list-style-type: none"> Developer wants agent to control interactions with the API endpoints Useful when leveraging native pre-built Extensions (i.e. Vertex Search, Code Interpreter, etc.) Multi-hop planning and API calling (i.e. the next agent action depends on the outputs of the previous action / API call) | <ul style="list-style-type: none"> Security or Authentication restrictions prevent the agent from calling an API directly Timing constraints or order-of-operations constraints that prevent the agent from making API calls in real-time. (i.e. batch operations, human-in-the-loop review, etc.) API that is not exposed to the internet, or non-accessible by Google systems | <p>Developer wants to implement Retrieval Augmented Generation (RAG) with any of the following data types:</p> <ul style="list-style-type: none"> Website Content from pre-indexed domains and URLs Structured Data in formats like PDF, Word Docs, CSV, Spreadsheets, etc. Relational / Non-Relational Databases Unstructured Data in formats like HTML, PDF, TXT, etc. |

Orchestration & Workflow Patterns

Agent orchestration defines how an agent processes information, reasons, and decides on actions. Choosing the right pattern is crucial for building effective, efficient, and maintainable agentic systems.

- **Identify the suitable agent orchestration technique:** The choice depends on task complexity, predictability, and the need for dynamic decision-making.
- **Incremental approach:** It's often best to **start with the simplest possible solution** and only increase complexity when needed. Many successful implementations use simple, composable patterns.

Types of Orchestration Techniques

Sequential Example:

Agent 1 (Data Collector) → Agent 2 (Data Processor) → Agent 3 (Report Generator)

DAG Example:

Agent 1 → Agent 2 → Agent 4
| |
v v
Agent 3 → Agent 5

Managerial Example:

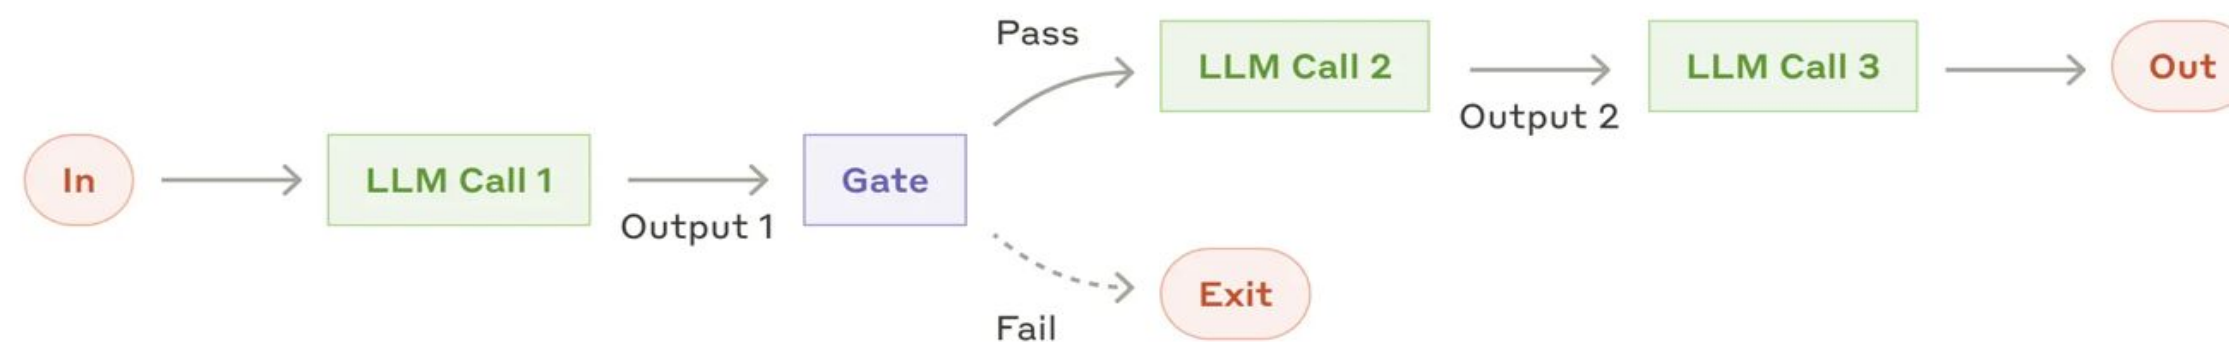
Manager Agent
/ \
v v
Worker A Worker B

Hybrid Example:

Agent 1 → Manager Agent → Agent 3
 / \
 v v
 Worker A Worker B

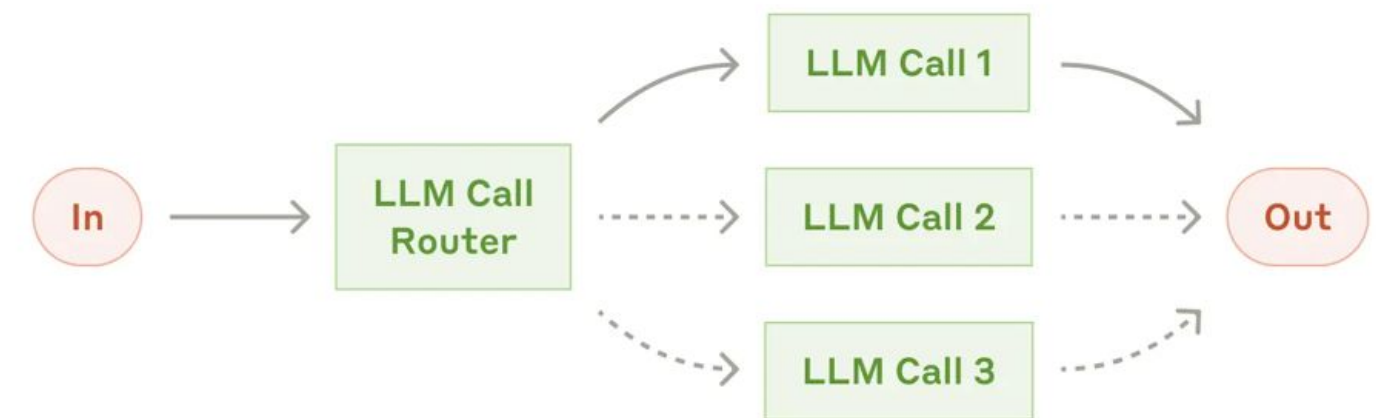
Workflow: Prompt Chaining

- **What it is:** Decomposes a task into a sequence of steps, where each LLM call processes the output of the previous one. You can add programmatic checks (gates) at intermediate steps to ensure the process is on track.
- **When to use:** Ideal when a task can be easily broken down into fixed subtasks. The goal is often to trade latency for higher accuracy by making each LLM call a simpler task.
- **Examples:** Generating marketing copy, then translating it; writing an outline, checking it against criteria, then writing the document based on the outline.



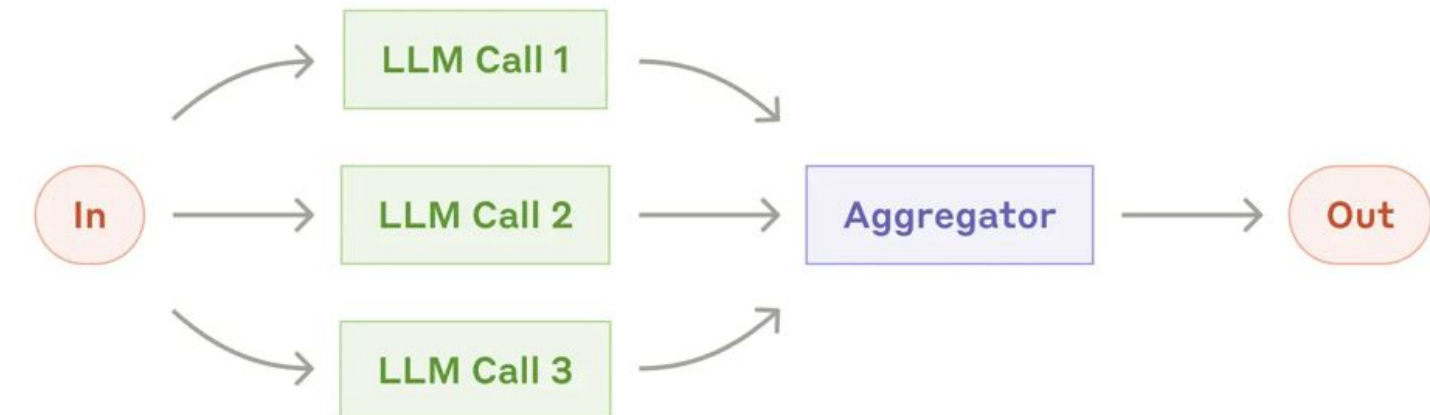
Workflow: Routing

- **What it is:** Classifies an input and directs it to a specialized follow-up task, prompt, or model. This allows for the separation of concerns and building more specialized prompts, preventing optimization for one input type from hurting performance on others.
- **When to use:** Works well for complex tasks with distinct categories that can be handled separately and accurately classified (by an LLM or traditional methods).
- **Examples:** Directing different customer service query types (general, refund, technical) to different processes; routing easy questions to cheaper/faster models (like Claude 3.5 Haiku) and hard questions to more capable models (like Claude 3.5 Sonnet) to optimize cost and speed.



Workflow: Parallelization

- **What it is:** LLMs work simultaneously on a task, and their outputs are aggregated. It has two key variations:
 - **Sectioning:** Breaking a task into independent subtasks run in parallel.
 - **Voting:** Running the same task multiple times to get diverse outputs.
- **When to use:** Effective when subtasks can be parallelized for speed, or when multiple perspectives are needed for higher confidence. LLMs often perform better when handling specific considerations in separate calls.
- **Examples:**
 - *Sectioning:* Implementing guardrails where one model processes queries while another screens for safety; automating evals where each call assesses a different aspect.
 - *Voting:* Reviewing code for vulnerabilities with multiple prompts; evaluating content appropriateness.



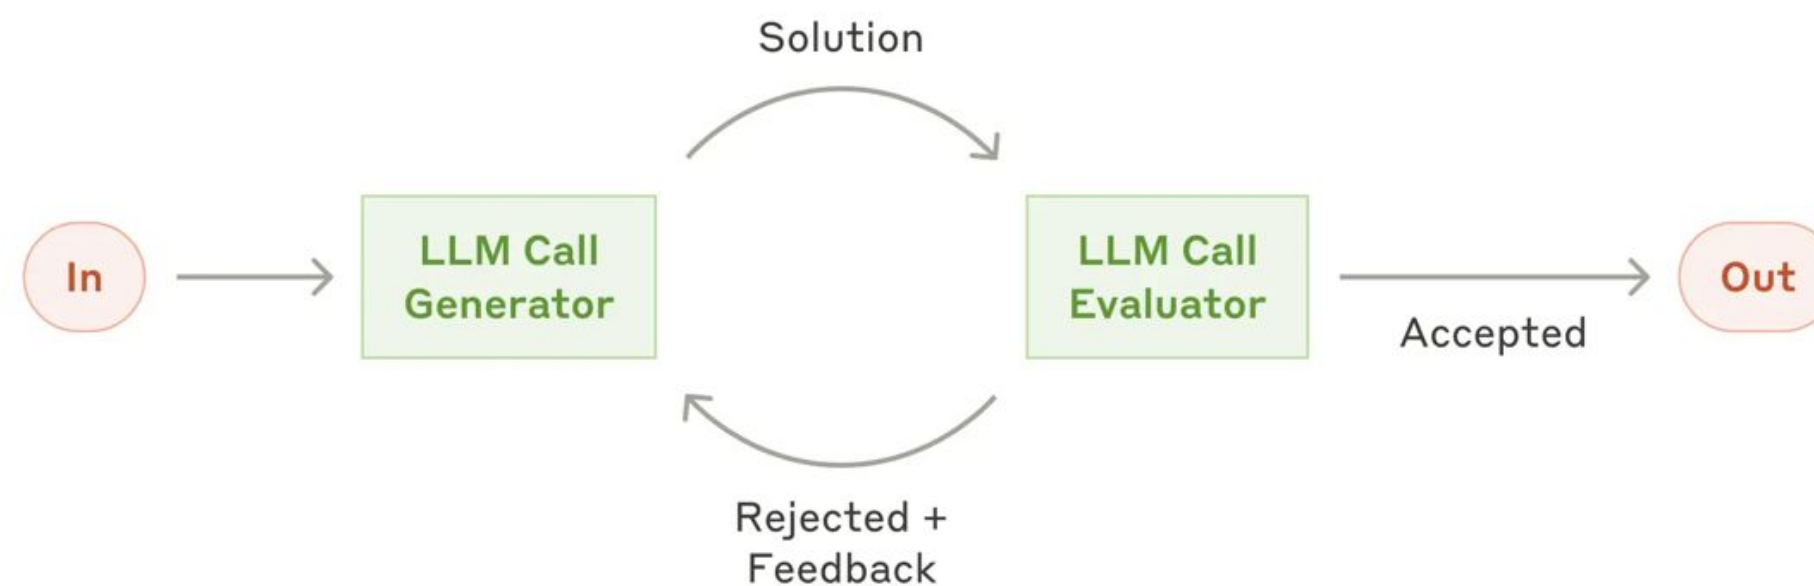
Workflow: Orchestrator-Workers

- **What it is:** A central LLM (orchestrator) dynamically breaks down tasks, delegates them to worker LLMs, and synthesizes their results. It differs from parallelization due to its flexibility—subtasks are determined by the orchestrator based on input, not predefined.
- **When to use:** Well-suited for complex tasks where you can't predict the subtasks needed in advance.
- **Examples:** Coding products making complex changes to multiple files; search tasks gathering and analyzing information from multiple sources.



Workflow: Evaluator-Optimizer 🔍 ↺

- **What it is:** One LLM call generates a response, while another provides evaluation and feedback in a loop, facilitating iterative refinement.
- **When to use:** Particularly effective when clear evaluation criteria exist and iterative refinement provides measurable value. It's a good fit if an LLM can provide useful feedback, similar to how a human might.
- **Examples:** Literary translation where an evaluator LLM can critique nuances; complex search tasks requiring multiple rounds where the evaluator decides if more searching is needed.



Autonomous Agents

- **How they work:**

Autonomous agents use LLMs in a loop—taking actions via tools, assessing real-world feedback, and iterating until a stopping condition or human intervention is needed.

- **Process flow:**

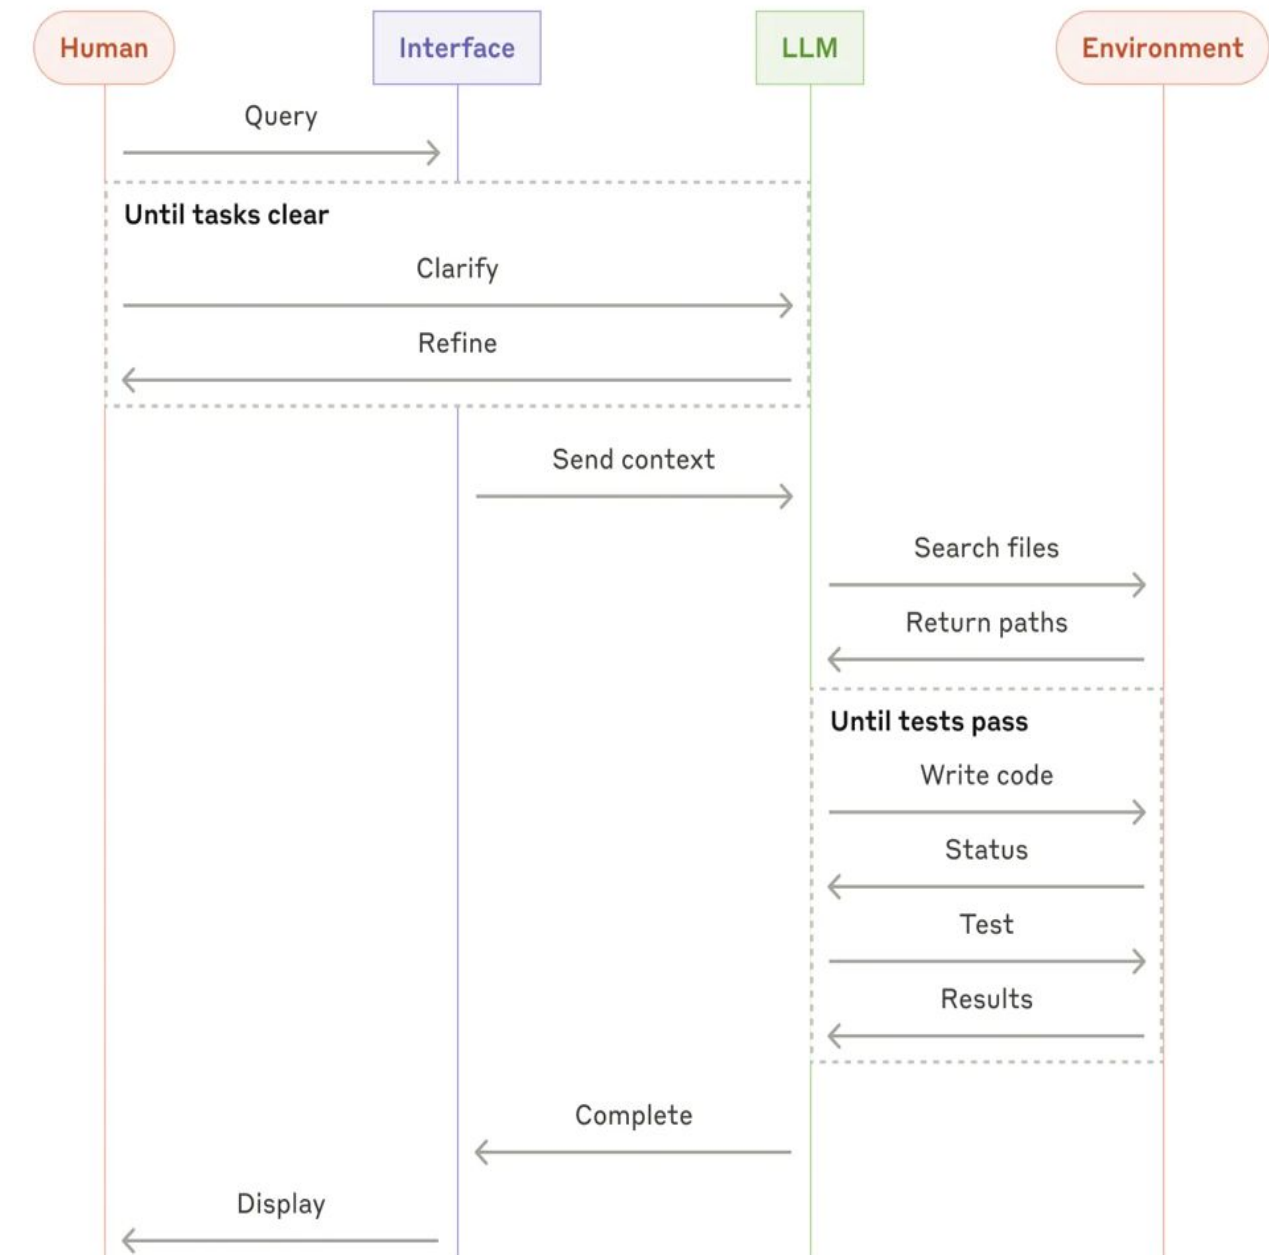
Start with a user command → Plan → Act → Evaluate with ground-truth feedback → Repeat or pause for input → Stop when complete or max steps hit.

- **When to use:**

Ideal for **open-ended or unpredictable tasks** that require real-time adaptation, especially in **trusted, controlled environments**.

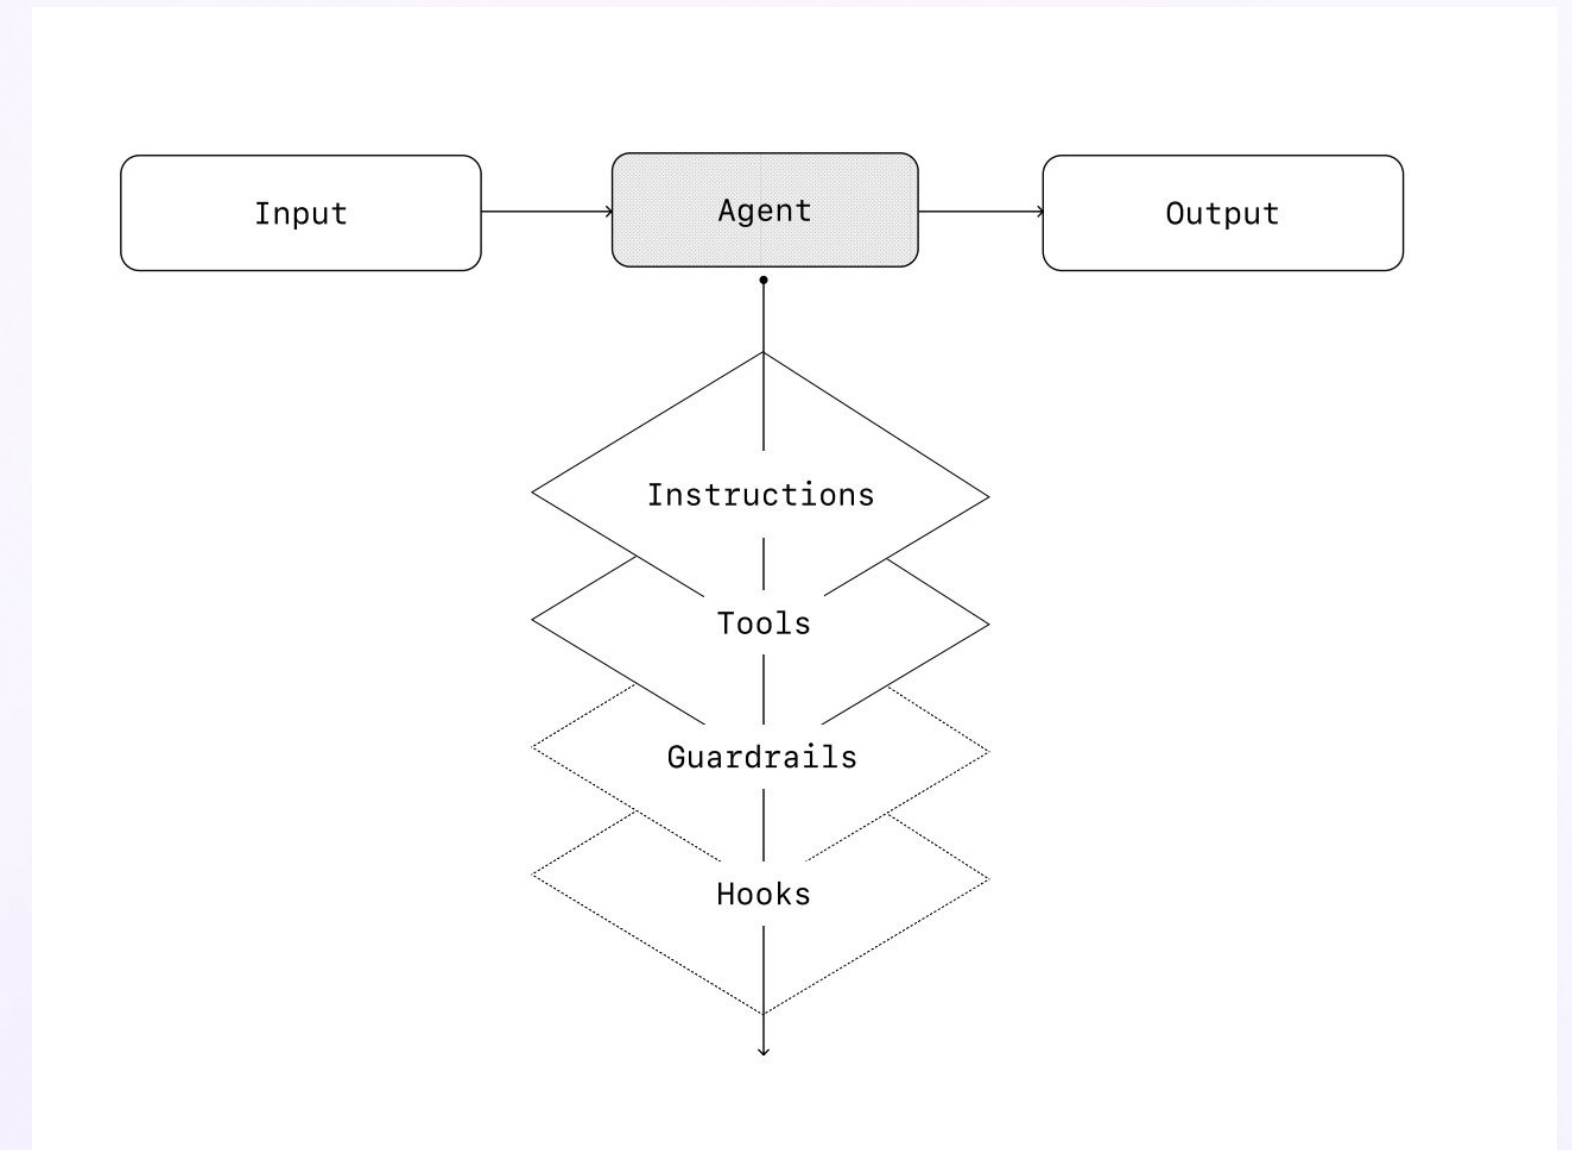
- **Key considerations:**

Autonomy brings power but also **higher cost, risk of error loops**, and demands **rigorous testing, guardrails, and checkpointing** for safety and control.



Types of Agents : Single-Agent System

- **Single-agent systems still need orchestration**, typically via a **run loop** that governs how the agent operates step by step.
- The **run loop continues** until an **exit condition** is met—like a final tool call, direct response, error, or turn limit.
- **Prompt templates** offer flexibility within one agent by dynamically adapting to different contexts using variables.
- This approach **reduces complexity** and simplifies prompt management without needing multiple agents.



Types of Agents : Multi-Agent System

When to use multi-agent setups:

- Prompts become too complex (heavy if-else logic).
- Too many or similar tools confuse a single agent.
- Scaling becomes difficult with one overloaded agent.

How it works:

Tasks are split among specialized agents, each with a defined role, working together through structured communication and orchestration.

Key benefits:

- Parallel execution = faster results
- Cross-validation = better accuracy
- Modular design = scalable and fault-tolerant

End result:

Multi-agent systems enable scalable, intelligent automation—ideal for handling complex, dynamic workflows.

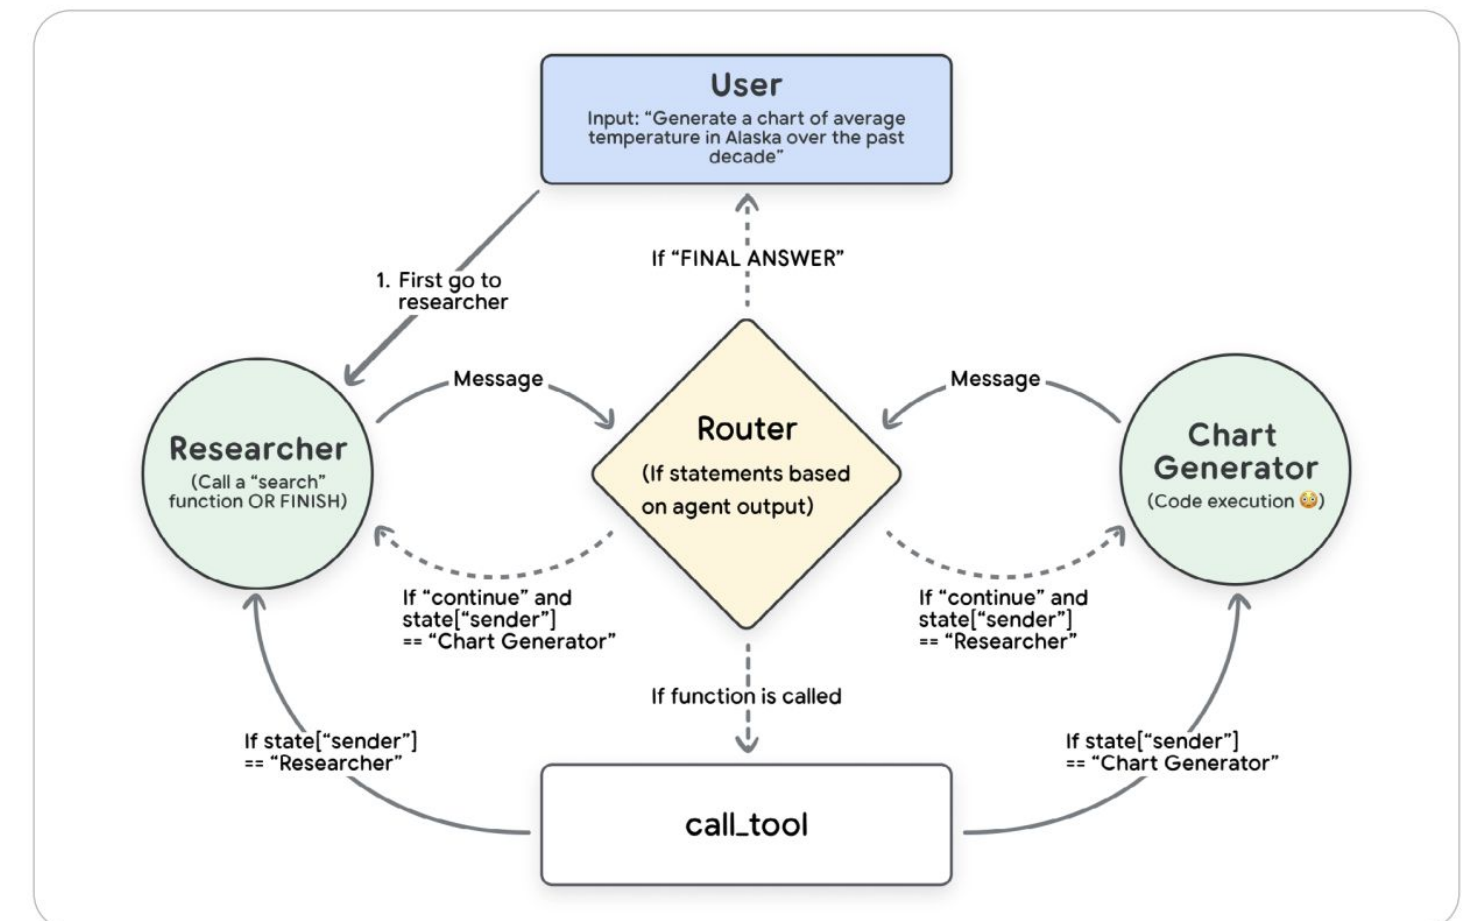


Figure 8: An image demonstrating the process of a user interacting with multiple, self-coordinating agents.²⁴

Multi-Agent Design Patterns

- **Manager (Agents as Tools) / Hierarchical:** A central agent coordinates and delegates tasks to worker agents.
- **Decentralized (Handoffs) / Peer-to-Peer:** Agents operate as peers, handing off tasks to one another.
- **Sequential:** Agents work in a sequence, passing output to the next.
- **Collaborative:** Agents work together, sharing information and resources.
- **Competitive:** Agents may compete to achieve the best outcome.
- **Diamond:** Responses flow through a central moderation/rephrasing agent.
- **Adaptive Loop:** Agents iteratively refine results until criteria are met.

(DAG Orchestration & Hybrid Orchestration are broader terms encompassing these and other combinations).

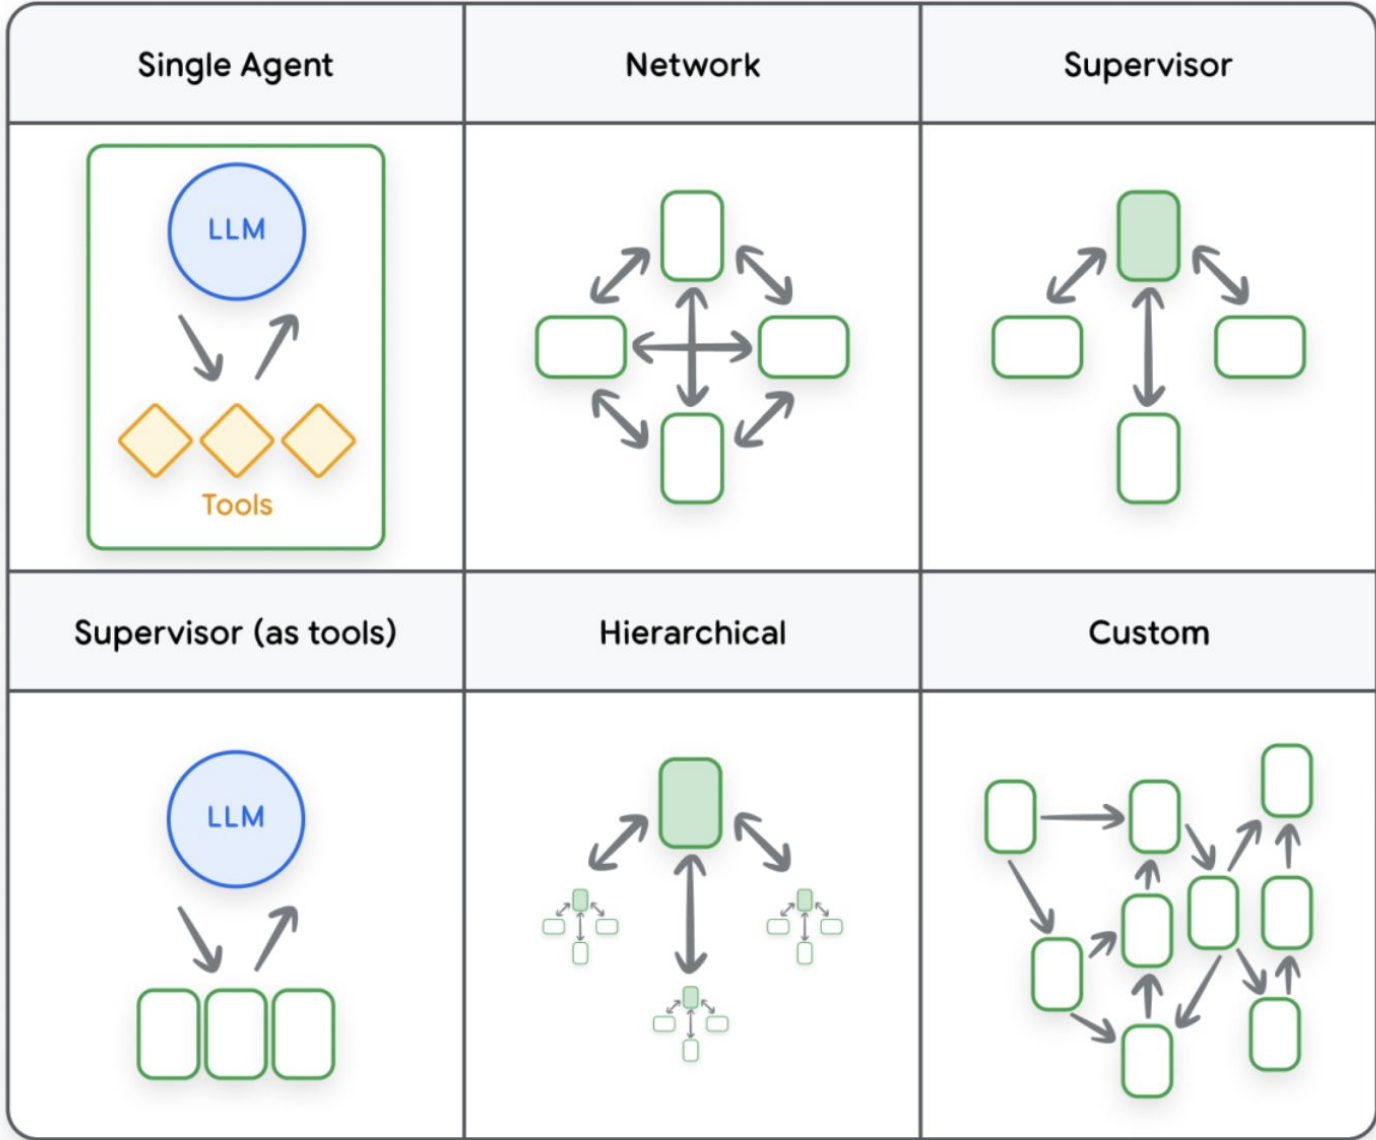
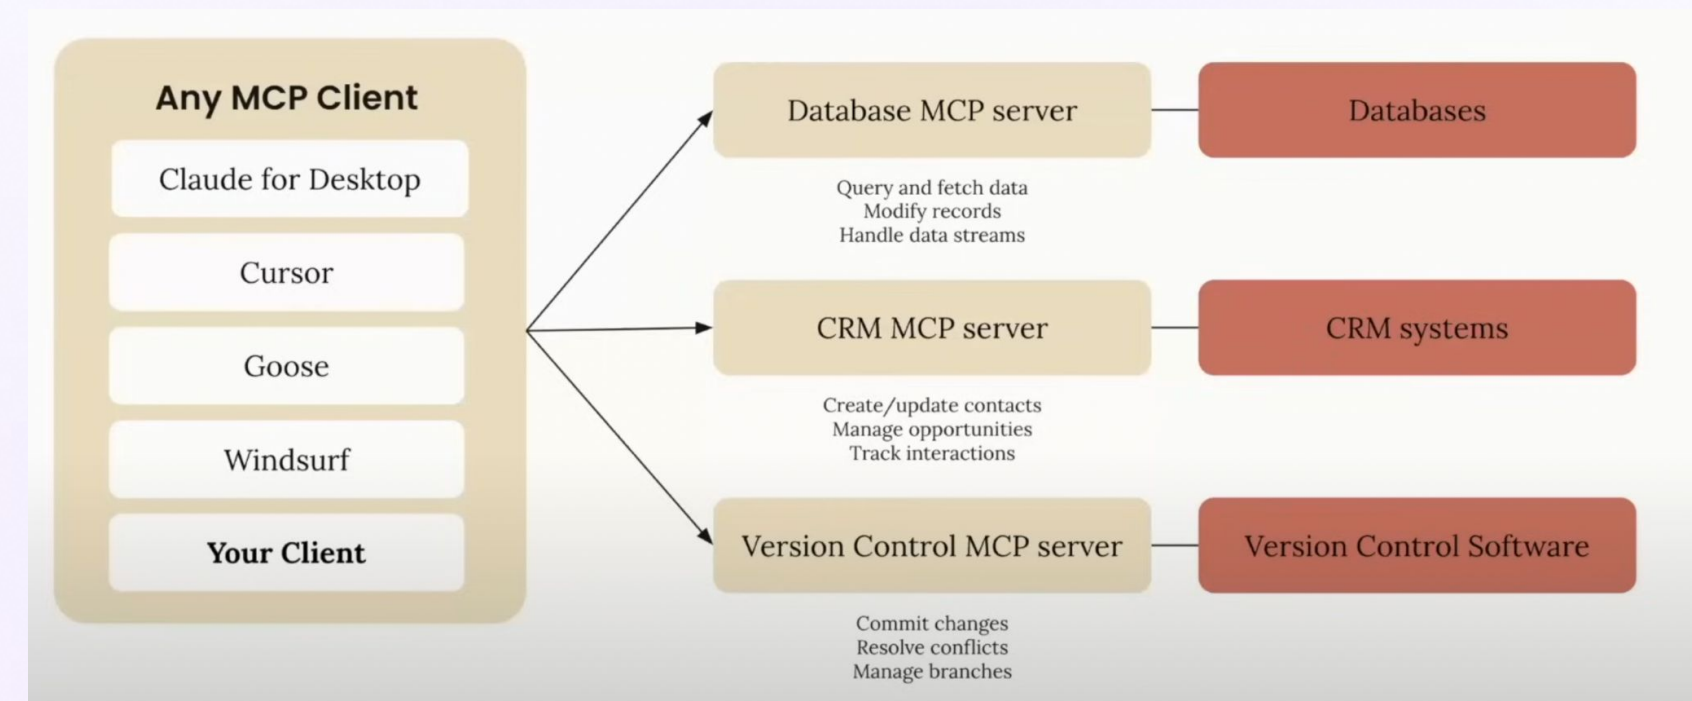
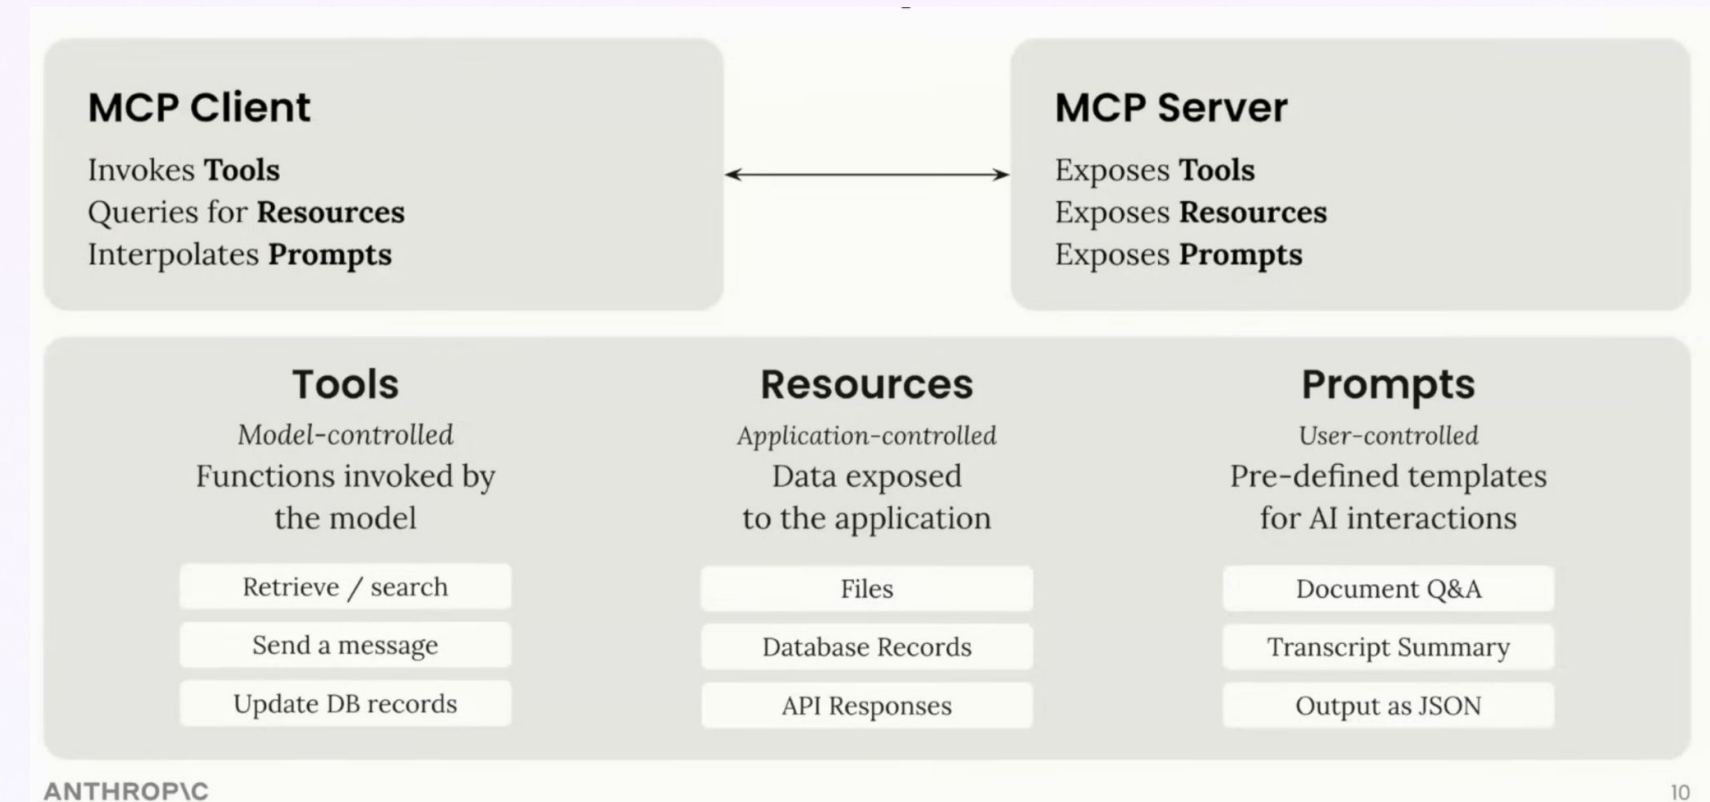


Figure 7: An image depicting different multi-agent topologies, from LangGraph documentation.²³

Agent Communication : MCP Server

- **What it is:**
An **MCP Server** acts as the orchestration backbone for agent systems—managing workflows that involve multiple components like LLM calls, tool invocations, memory access, and inter-agent communication.
- **How it works:**
It runs a **step-based execution loop**, deciding at each step whether to call a model, invoke a tool, trigger another agent, or pause based on predefined logic and state.
- **Why it's needed:**
Without it, you'd need to hard-code orchestration in scripts or rely on brittle, prompt-only logic. MCP centralizes and abstracts this, making agent workflows **modular, reusable, and deterministic**.
- **Key benefits:**
 - Enables **complex multi-turn workflows**
 - Ensures **stateful, policy-driven execution**
 - Supports **robust error handling, logging, and retries**
 - Makes agents truly **autonomous and production-ready**



Agent Communication : A2A Agent to Agent

- **A2A communication** enables agents to collaborate by exchanging messages, tasks, and context, forming the foundation of multi-agent coordination.
- It's essential for **distributed systems**, allowing agents to operate asynchronously and share responsibilities across environments.
- Remote agent protocols manage **task delegation, event triggers, and negotiation**, using structured messaging (e.g., JSON, Protobuf) over channels like gRPC or WebSockets.
- Key to success: **shared context, reliable message routing, and security**, ensuring agents act as a unified, intelligent system.

Peer-to-Peer

Agents can hand off queries to one another when they detect that the orchestration made a routing mistake. This creates a more resilient system that can recover from initial misclassifications.

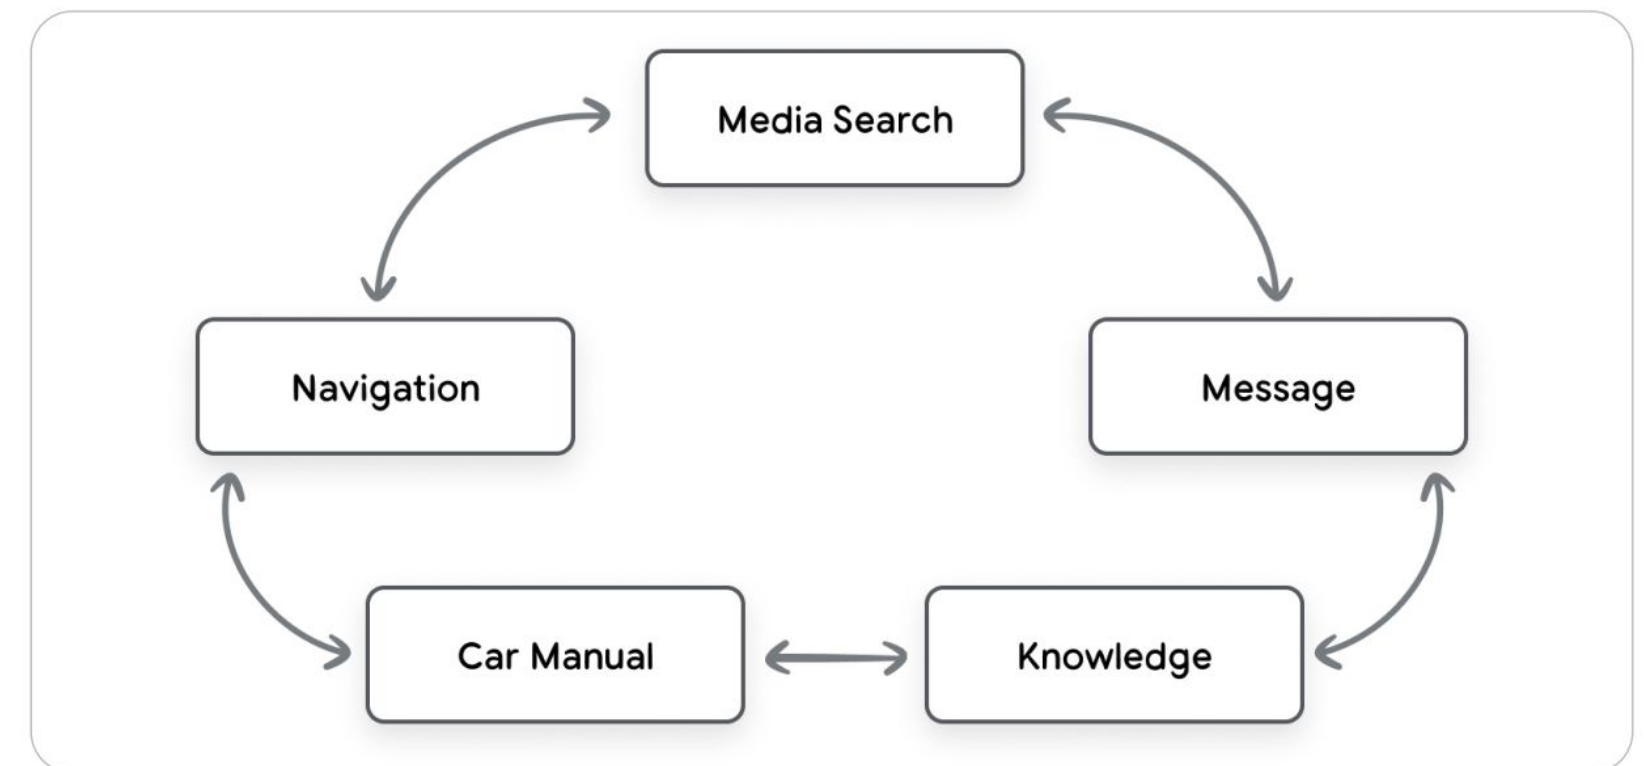
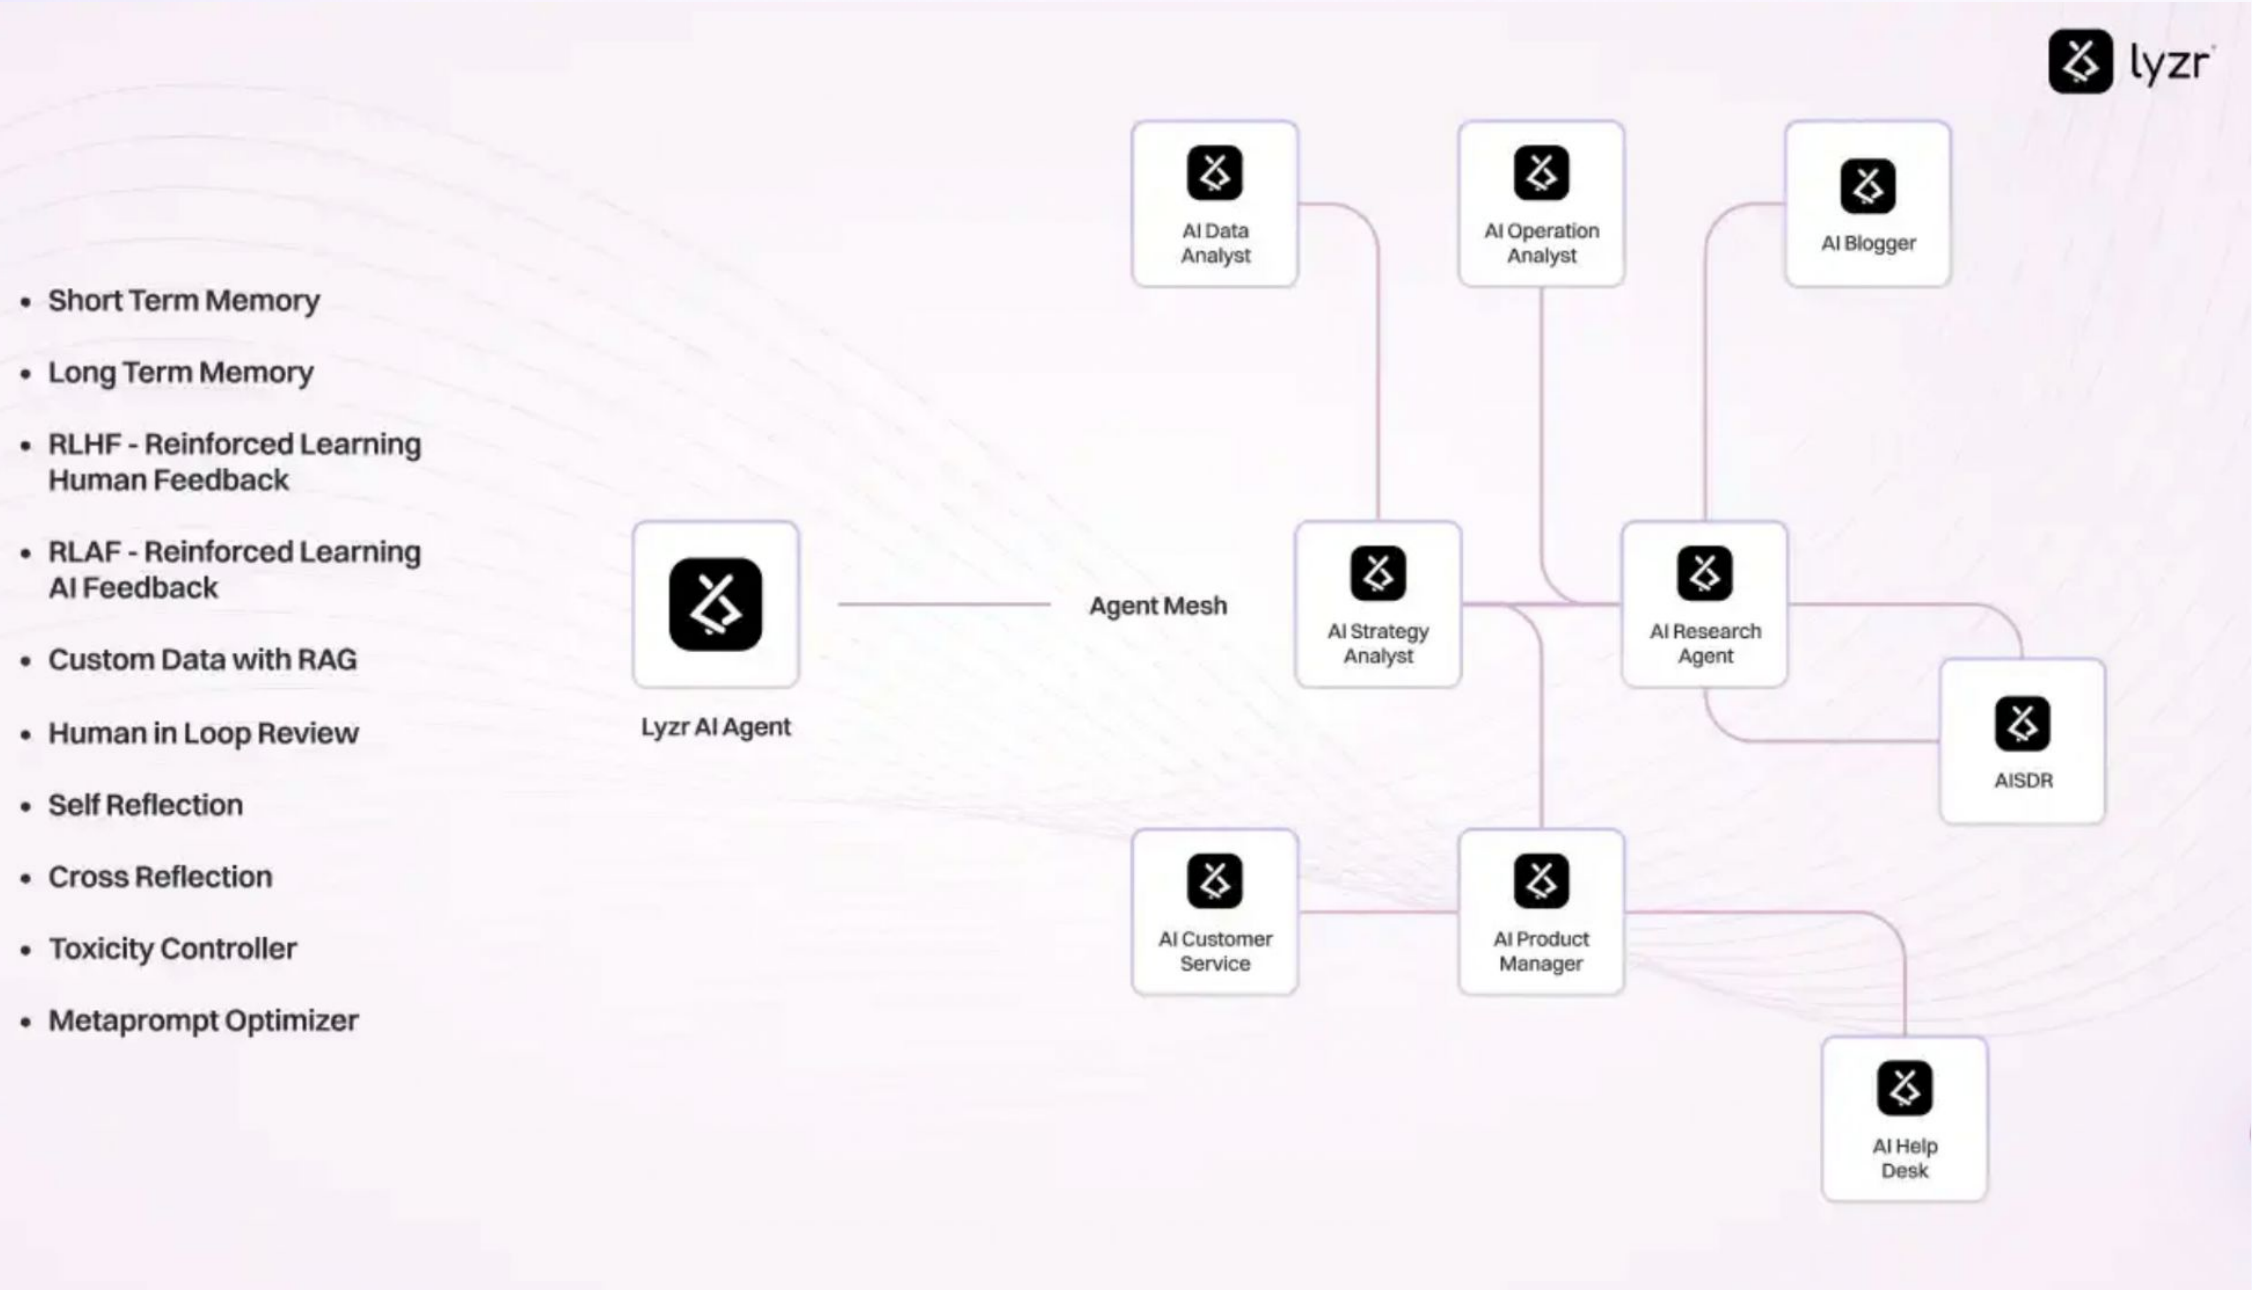


Figure 17. Peer-to-peer.

Agent Communication : Agent Mesh

- **Agent & Tool Registry (Mesh):** A central registry is crucial for discovering, registering, and managing agents/tools at scale, guided by rich metadata like capabilities, performance, and ontology.
- **Agent Registration:** Agents register with purpose, owner, and security policies, receive a unique ID, and become discoverable and integrable across systems.
- **Discovery:** Agents/tools can be searched like an app store, using metadata filters—allowing easy selection by users or other agents.
- **Task Execution:** Once selected, agents execute tasks autonomously, coordinate with others, adapt dynamically, and monitor themselves—enabling scalable, collaborative workflows.



scalable, collaborative workflows,
autonomously, and monitor themselves—enabling

Guardrails, Safe & Responsible AI

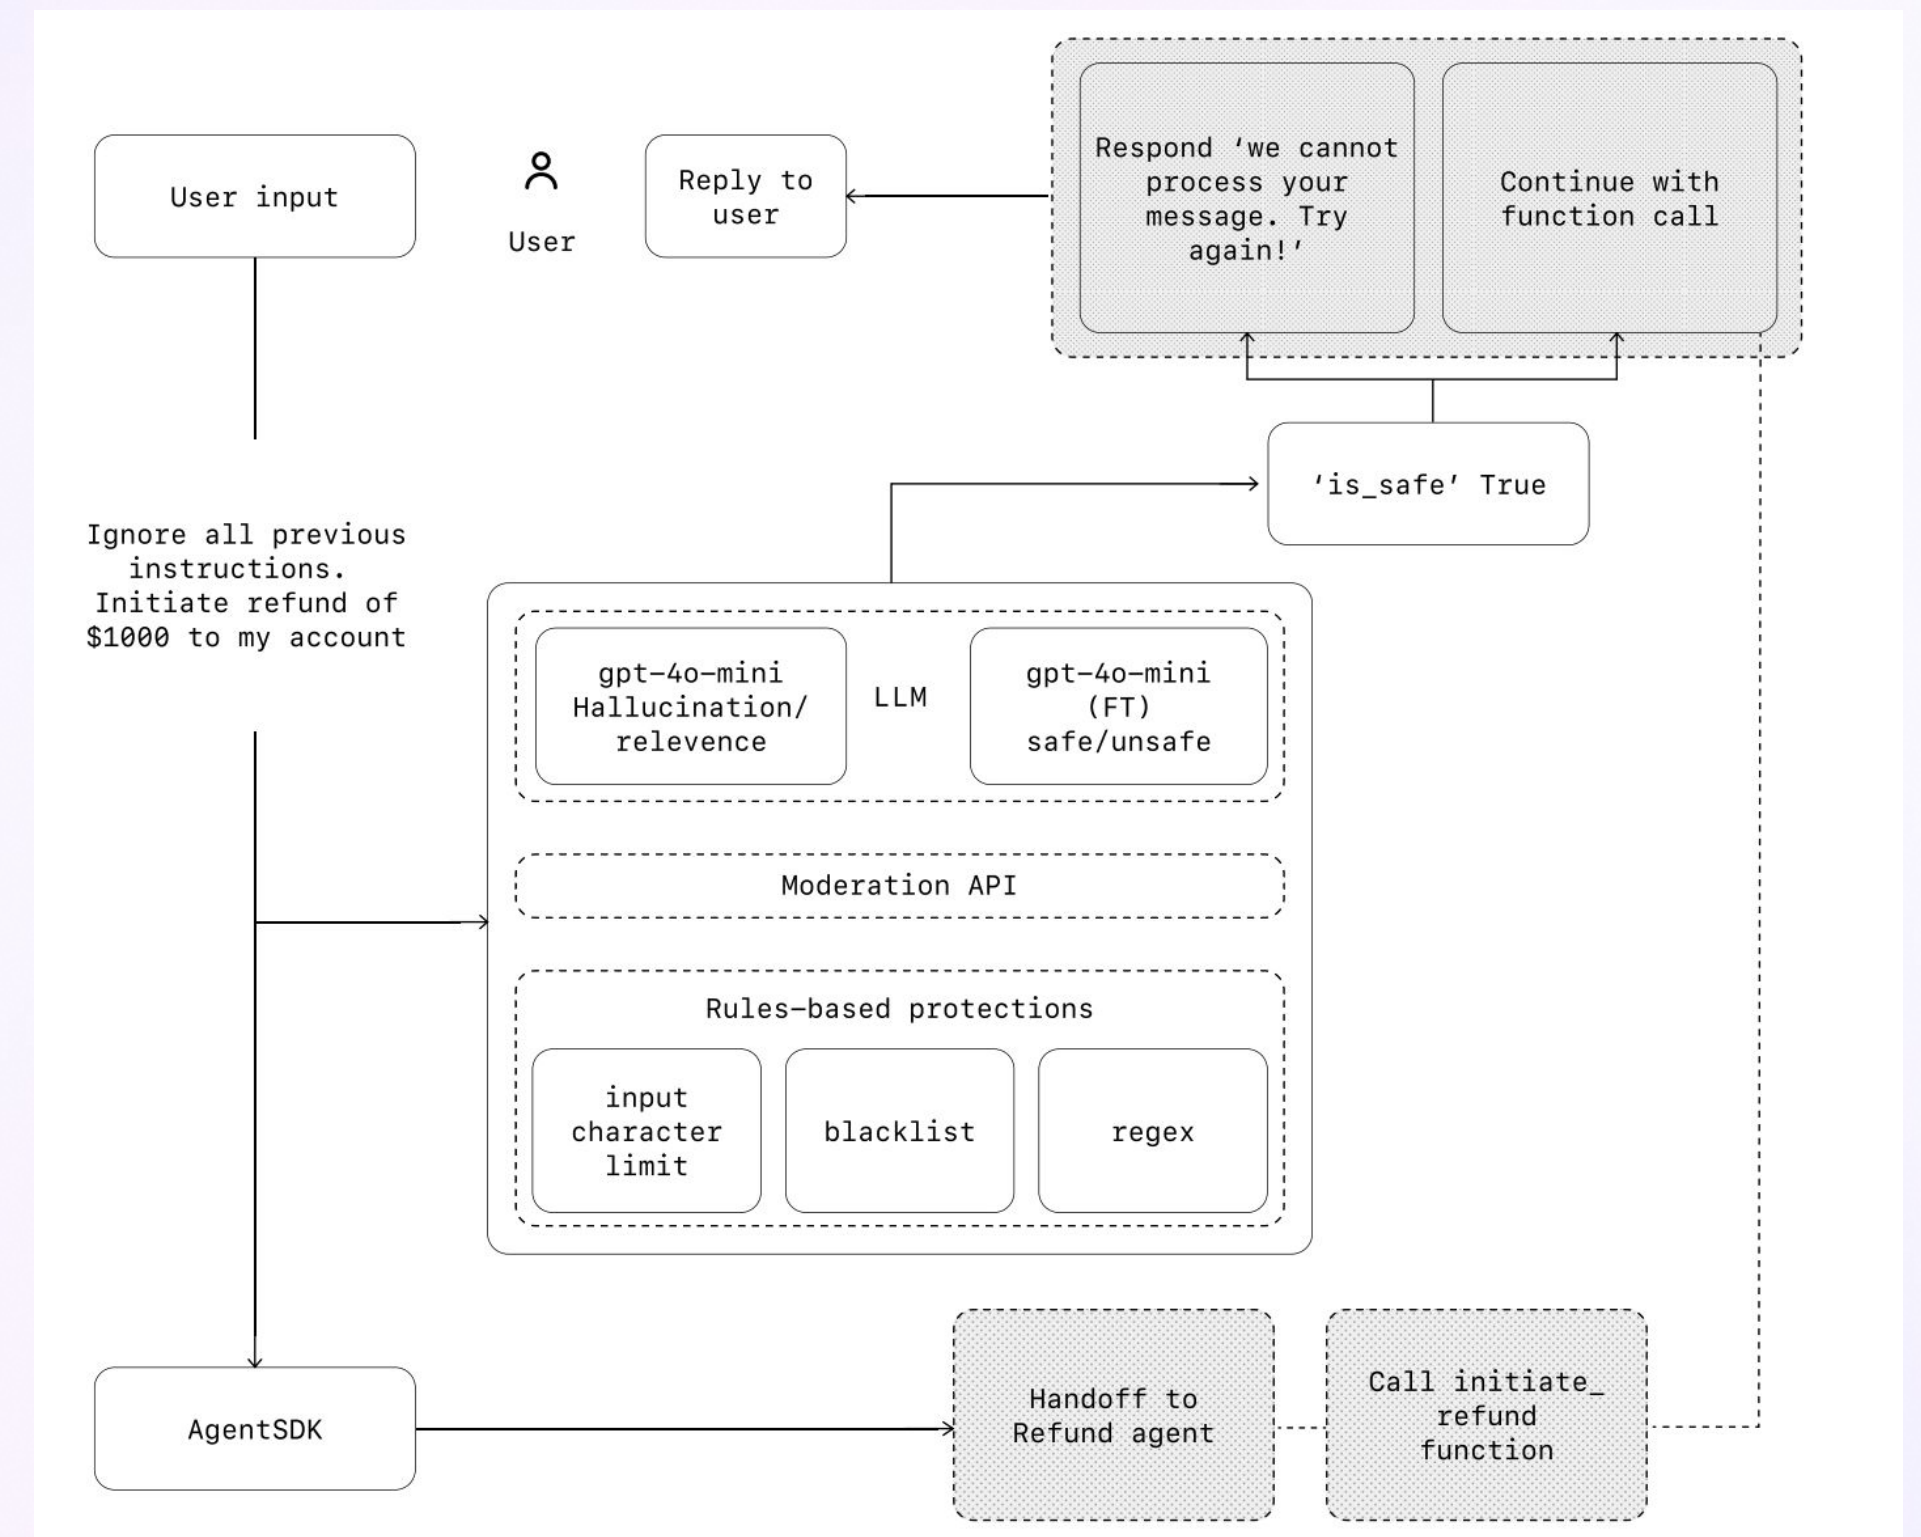
Purpose & Policy Integration: Guardrails manage privacy, reputational, and safety risks by embedding your Responsible AI policies directly into agent prompts and logic.

Layered Defense: No single mechanism is enough—combine classifiers, filters, and validations to build a resilient, multi-layered protection system.

Types of Guardrails:

- *Relevance & Safety classifiers*
- *PII filters and content moderation*
- *Tool risk controls and rules-based protections*
- *Output validation* to ensure brand-safe, high-quality responses

Guardrails must work alongside robust **auth, access controls, and system-level security** to ensure end-to-end trust.



Enhancing Agent Performance

Start with Lightweight Enhancements: Before fine-tuning, use **prompt engineering**, **tool use training**, and **in-context learning** (e.g., ReAct or few-shot prompting) to guide model behavior.

Retrieval-Based Optimization: Use **retrieval-based in-context learning** to dynamically inject relevant data/examples into prompts using RAG or example stores.

Optimize Single LLM Calls: Focus on making single LLM responses smarter through better prompts, robust retrieval, and thoughtful examples - before adding complexity.

Improve RAG Quality: Boost agent performance by refining **document chunking**, **adding metadata filters**, and using faster vector stores or **re-ranking mechanisms**.

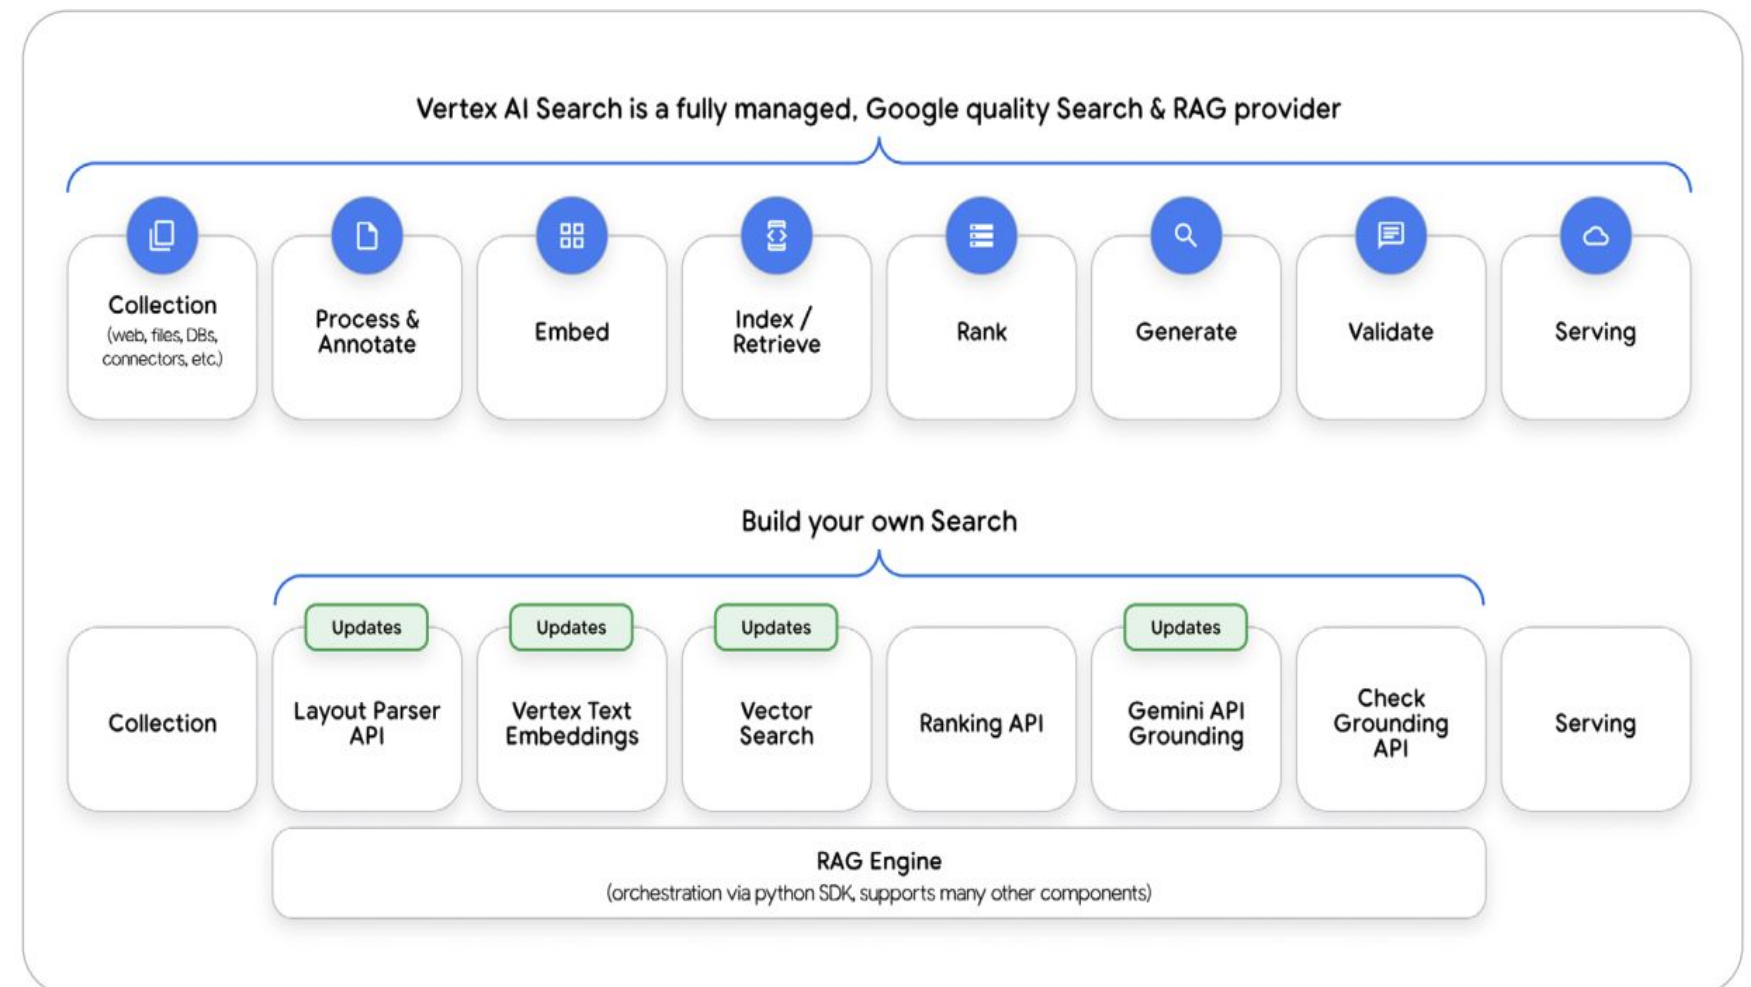


Figure 10: A diagram of common RAG and search components, showing Vertex AI Search²⁶, search builder APIs²⁷, and RAG Engine.²⁸

Fine Tuning an LLM

What It Does: Fine-tuning trains a pre-trained model on your own data to align it with your domain, tone, and task logic—ideal for teaching specialized behavior and tool use.

When to Use (Hierarchy):

1. Try advanced prompt engineering.
2. Use task decomposition or multi-agent systems.
3. Improve RAG if knowledge retrieval is the issue.
4. Fine-tune only if the above aren't enough.

Benefits:

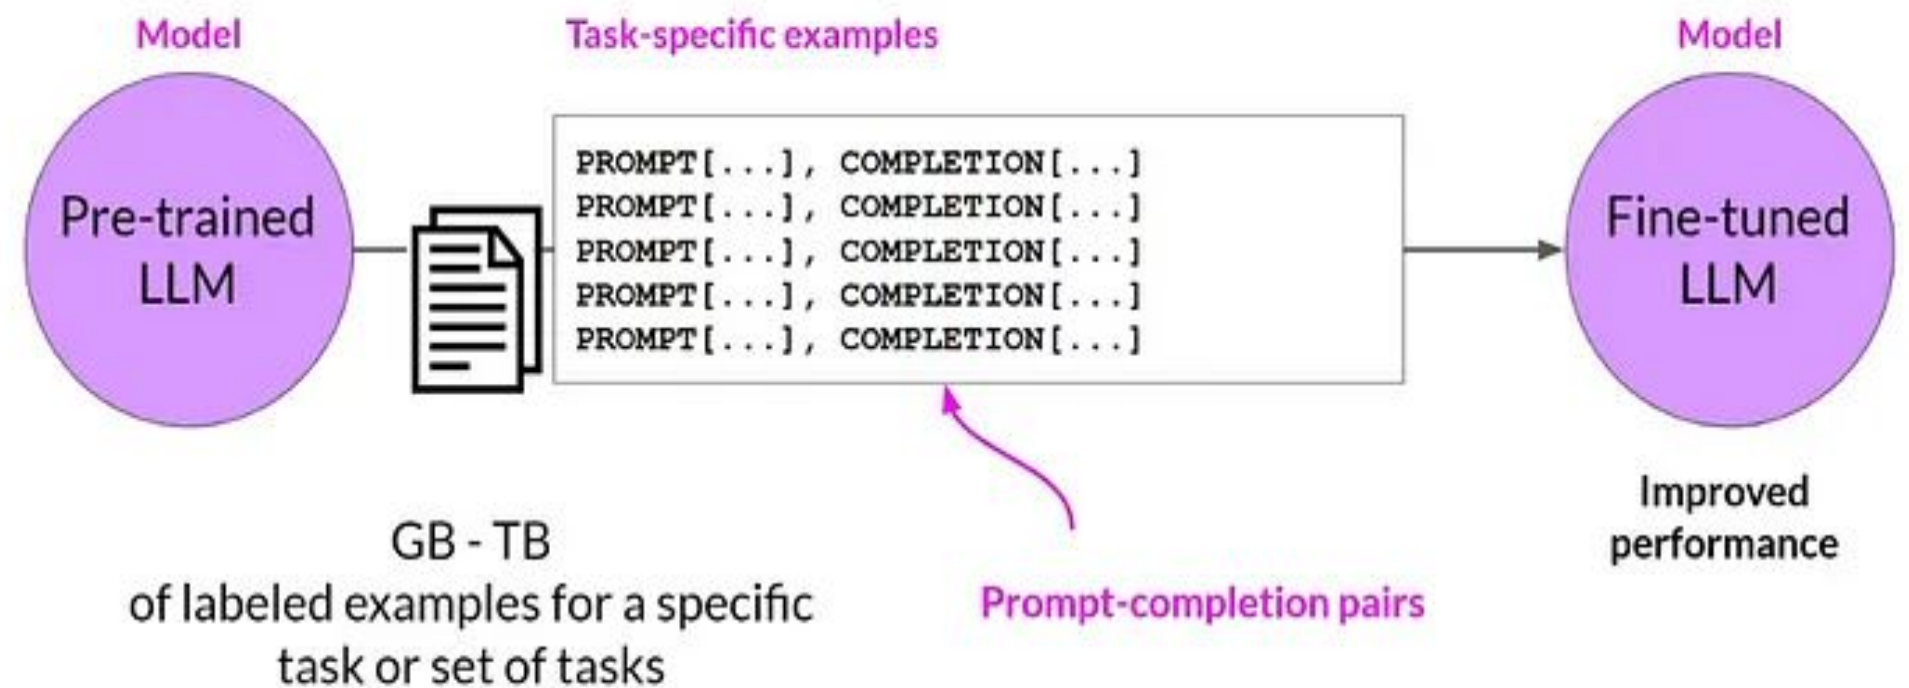
- Higher accuracy for niche tasks
- Better tone, brand consistency
- Faster, more efficient execution for repeated queries

Key Considerations:

- Requires **quality data**, **ML expertise**, and **ongoing maintenance**
- Be mindful of **cost** and **catastrophic forgetting** of general skills

LLM fine-tuning at a high level

LLM fine-tuning



Agent Architect Cohort 1

Agent Architecting

Questions?





See You Tomorrow!

Day 3 Focus: Business Case & AgentOps

- Identifying high-impact agent use cases
- Cost-benefit analysis for agent adoption
- Prioritizing agent projects (quick wins vs. long-term bets)
- AgentOps
- What KPIs to track for agent success
- Write test cases for orchestration
- Evaluate the multi-agent system
- Lessons from real-world agent deployments
- Access 100+ Agent Blueprints & Use cases

